

Online Optimization of Neural Network Controls for Stochastic Differential Equations



1077710
Kellogg College
University of Oxford

A thesis submitted for the degree of
the MSc in Mathematical Finance

June 28, 2024

Acknowledgements

I would like to express my gratitude to my supervisor, Justin Sirignano, for his invaluable guidance and encouragement throughout this dissertation.

I extend my appreciation to my teachers in the Department of Math, for their wonderful teaching and assistance.

On a personal note, I would like to thank my parents, for their endless support and belief in me.

Abstract

This dissertation focuses on the online optimization of neural network controls for stochastic differential equations (SDEs), addressing the computational challenges posed by high-dimensional stochastic optimal control problems. Traditional numerical methods become infeasible in high-dimensional settings, necessitating the development of novel online optimization algorithms. The dissertation re-implements examples from previous works and extends the research to higher dimensions, correlated processes, and nonlinear cases. The core contribution is the application of continuous-time stochastic gradient descent for training neural network controls, validated through mathematical proofs and extensive numerical experiments. The results demonstrate the algorithm's effectiveness in achieving optimal control, even in complex and high-dimensional scenarios, with potential applications in mathematical finance and engineering.

Contents

1	Introduction	1
1.1	Backgrounds	1
1.2	Objectives of This Dissertation	1
1.3	Paper Outline	2
1.4	Summary of The Numerical Results	2
2	Main Theorem	3
2.1	The Main Theorem	3
2.2	Intuitive Inspiration	4
2.2.1	The Natural Approach	4
2.2.2	The New Approach	5
2.3	Summary of the Mathematical Proof	6
2.3.1	Proposition	6
2.3.2	Theorem 2.2 of the original paper [4]	7
2.4	Computational Limitations	7
2.4.1	Dependence on Hyper-parameters	7
2.4.2	Stochastic Estimates and Noise	8
2.4.3	Convergence to Global Minima	8
2.4.4	High-dimensional Problems	8
2.4.5	Path-dependent and Nonlinear SDEs	8
2.4.6	Application-specific Challenges	8
3	Implementation	9
3.1	Simple Cases	9
3.1.1	One-Dimensional Ornstein-Uhlenbeck Process	9
3.1.2	One-Dimensional Nonlinear Process	14
3.1.3	Optimizing over the Drift and Volatility Coefficients	15
3.1.4	Multi-Dimensional Ornstein-Uhlenbeck Process	18

3.1.5	Path-dependent SDE	22
3.2	Linear-quadratic Regulator	23
3.2.1	One-dimensional Linear Control	24
3.2.2	Multi-dimensional Linear Control	25
3.3	Neural Network Control	27
3.3.1	One-dimensional Neural Network Control	27
3.3.2	Multi-dimensional Neural Network Control	29
4	Futher Research	32
4.1	Implementation Method Innovation	32
4.1.1	Methodology Given In The Original Paper	32
4.1.2	An Easier Way	32
4.1.3	Advantages	34
4.2	Higher Dimensions	35
4.2.1	LQR	35
4.2.2	Neural Network Control	35
4.2.3	Conclusions	36
4.3	Correlated OU Processes	36
5	Conclusions	39
5.1	Summary of Main Findings	39
5.2	Future Research	39
A	Key Code	41
A.1	One-dimensional Neural Network Control	41
A.2	Multi-dimensional Neural Network Control	47
A.3	LQR	54
Bibliography		58

List of Figures

3.1	$J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y - 2)^2$	9
3.2	#epochs = 100	10
3.3	$J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y^2 - 2)^2$	11
3.4	Two Parameter case	12
3.5	Non linear case	15
3.6	Drift and Volatility Linear	16
3.7	Drift and Volatility Linear	17
3.8	Drift and Volatility Non Linear	18
3.9	Independent m=3	19
3.10	Independent m=10	19
3.11	Correlated m=3	21
3.12	Correlated m=3 refined	21
3.13	Correlated m=10 refined	22
3.14	Path Dependent	23
3.15	LQR 1-D	25
3.16	LQR Multi-D n=5	27
3.17	LQR Multi-D n=20	27
3.18	NN control 1-D	29
3.19	NN control 1-D Gradient Check	29
3.20	NN control Multi-D n=5	31
4.1	Hyper Cube	34
4.2	LQR n=50	35
4.3	NN n=30	36
4.4	Correlated dim=10 (1)	37
4.5	Correlated dim=10 (2)	38

Chapter 1

Introduction

1.1 Backgrounds

Stochastic optimal control is a crucial field in applied mathematics and engineering. It is vital due to its ability to address uncertainty in dynamic systems, improve decision-making, drive mathematical innovation, and impact a wide range of disciplines and applications. Such as mathematical finance. For example, in Portfolio Optimisation and Optimal Execution problems.

The problems we are facing is that, high-dimensional stochastic optimal control presents significant computational challenges. Since the optimal control must satisfy the Hamilton-Jacobi-Bellman (HJB) equation, which becomes computationally infeasible to solve in high-dimensional settings. Hence we need to implement a new on-line optimization method for training neural network controls for SDEs that remains computationally feasible in high dimensions. The paper Continuous-time stochastic gradient descent for optimizing over the stationary distribution of stochastic differential equations, by Wang, Ziheng and Sirignano, Justin [4] have given a successful implementation of the method which is feasible in the high dimensional cases.

1.2 Objectives of This Dissertation

In this dissertation we re-implemented those examples and conducted research beyond the original paper, such as exploring higher dimensions, adding correlations, and studying cases of non-linear processes. Additionally, we went over the principle of this given method (algorithm) and verify the principles of the original paper.

1.3 Paper Outline

The thesis is structured as follows. First we explain the main theorem of the original paper and justify it. Then we use the given method as a tenet for the different algorithms for different cases. We do the implementation from simple to complicated cases. Next, in Chapter 4 we go beyond the original research to try some new things. Then, in the last chapter, we conclude and summarize the main findings of your research, as well as suggest directions for future research based on the findings and limitations.

1.4 Summary of The Numerical Results

The numerical results across various scenarios and dimensions demonstrate the efficacy of the online optimization algorithm in solving high-dimensional stochastic control problems. The algorithm shows good convergence properties, even in complex and high-dimensional cases, making it a valuable tool for applications in mathematical finance and engineering.

Chapter 2

Main Theorem

2.1 The Main Theorem

In the original paper[4]. Given a multi-dimensional Ornstein-Uhlenbeck process:

$$dX_t^\theta = (g(\theta) - h(\theta)X_t^\theta)dt + dW_t, \quad X_0^\theta = x \quad (2.1)$$

where $\theta \in \mathbb{R}^l$, $g(\theta) \in \mathbb{R}^d$, $h(\theta) \in \mathbb{R}^{d \times d}$, $W_t \in \mathbb{R}^d$, $X_t^\theta \in \mathbb{R}^d$ and σ is a scalar constant. Let π_θ be the stationary distribution of X_t^θ . Our goal is to solve the optimization problem

$$\min_{\theta} J(\theta) = \min_{\theta} (\mathbb{E}_{\pi_\theta} f(Y) - \beta)^2 \quad (2.2)$$

where β is a constant. To solve (2.2), we have the main online algorithm:

$$\begin{aligned} \frac{d\theta_t}{dt} &= -2\alpha_t (f(\bar{X}_t) - \beta) \nabla f(X_t) \tilde{X}_t \\ dX_t &= (g(\theta_t) - h(\theta_t)X_t) dt + \sigma dW_t \\ \frac{d\tilde{X}_t}{dt} &= \nabla_\theta g(\theta_t) - \nabla_\theta h(\theta_t) X_t - h(\theta_t) \tilde{X}_t \\ d\bar{X}_t &= (g(\theta_t) - h(\theta_t) \bar{X}_t) dt + \sigma d\bar{W}_t \end{aligned} \quad (2.3)$$

where W_t and \tilde{W}_t are independent Brownian motions, $\nabla_\theta g(\theta) \in \mathbb{R}^{d \times l}$, $\nabla_\theta h(\theta) \in \mathbb{R}^{d \times d \times l}$ and $\tilde{X}_t \in \mathbb{R}^{d \times l}$ is the gradient process for X_t .

To make our convergence theorem hold, we will require the following assumptions.

Assumption

- $g(\theta), \nabla_\theta^i g(\theta), h(\theta)$ and $\nabla_\theta^i h(\theta)$ are uniformly bounded functions for $i = 1, 2$.

- h is symmetric and uniformly positive definite, i.e. there exists a constant $c > 0$ such that

$$\min \{x^\top h(\theta)x\} \geq c|x|^2, \quad \forall \theta \in \mathbb{R}^\ell, x \in \mathbb{R}^d$$

- $f, \nabla^i f, i = 1, 2, 3$ are polynomially bounded.

$$|f(x)| + \sum_{i=1}^3 |\nabla^i f(x)| \leq C(1 + |x|^{\hat{m}}), \quad \forall x \in \mathbb{R}^d$$

for some constant $C, \hat{m} > 0$.

- The learning rate α_t satisfies $\int_0^\infty \alpha_t dt = \infty, \int_0^\infty \alpha_t^2 dt < \infty, \int_0^\infty |\alpha'_s| ds < \infty$, and there is a $\hat{p} > 0$ such that $\lim_{t \rightarrow \infty} \alpha_t^2 t^{\frac{1}{2}+2\hat{p}} = 0$.

Under these assumptions, we are able to give the following convergence result.

The Main Theorem

Under above assumptions and for the Ornstein-Uhlenbeck process (2.1), the algorithm (2.3) will converge to a stationary point almost surely:

$$\lim_{t \rightarrow \infty} |\nabla_\theta J(\theta_t)| \stackrel{\text{a.s.}}{=} 0 \quad (2.4)$$

2.2 Intuitive Inspiration

2.2.1 The Natural Approach

In the algorithm we use $(f(\bar{X}_t) - \beta) (\nabla f(X_t) \tilde{X}_t)^\top$ as a stochastic estimate for $\nabla_\theta J(\theta_t)$, where $J(\theta) = \sum_{n=1}^N (\mathbf{E}_{\pi_\theta}[f_n(Y)] - \beta_n)^2$, and it is used to update θ_t .

To better understand the main online algorithm (2.3), we first re-write the gradient of the objective function using the ergodicity of X_t^θ in the case of $N = 1$:

$$\begin{aligned} \nabla_\theta J(\theta) &= 2(\mathbf{E}_{Y \sim \pi_\theta} f(Y) - \beta) \nabla_\theta \mathbf{E}_{Y \sim \pi_\theta} f(Y) \\ &\stackrel{\text{a.s.}}{=} 2 \left(\lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T f(X_t^\theta) dt - \beta \right) \cdot \nabla_\theta \left(\lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T f(X_t^\theta) dt \right) \end{aligned}$$

Assume we have the requirements fulfilled, and the derivative and the limit can be interchanged, the gradient can be expressed as

$$\nabla_{\theta} J(\theta) = 2 \left(\lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T f(X_t^\theta) dt - \beta \right) \cdot \lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T \nabla f(X_t^\theta) \nabla_{\theta} X_t^\theta dt$$

Define $\tilde{X}_t^\theta = \nabla_{\theta} X_t^\theta$ and, under mild regularity conditions for the coefficients[3][2], \tilde{X}_t^θ will satisfy

$$d\tilde{X}_t^\theta = \left(\nabla_x \mu(X_t^\theta, \theta) \tilde{X}_t^\theta + \nabla_{\theta} \mu(X_t^\theta, \theta) \right) dt + \left(\nabla_x \sigma(X_t^\theta, \theta) \tilde{X}_t^\theta + \nabla_{\theta} \sigma(X_t^\theta, \theta) \right) dW_t$$

Note that \tilde{X}_t and \tilde{X}_t^θ satisfy the same equations, except θ is a fixed constant for \tilde{X}_t^θ while θ_t is updated continuously in time for \tilde{X}_t . Then, we have that

$$\nabla_{\theta} J(\theta) = 2 \left(\lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T f(X_t^\theta) dt - \beta \right) \cdot \lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T \nabla f(X_t^\theta) \tilde{X}_t^\theta dt$$

The formula above can be used to evaluate $\nabla_{\theta} J(\theta)$ and thus allows for optimization via a gradient descent algorithm. However, the problem is, X_t^θ must be given by simulation for a large time period $[0, T]$ for each optimization iteration (epoch), which is computationally costly.

2.2.2 The New Approach

An alternative way is to develop a continuous-time stochastic gradient descent algorithm which updates θ using a stochastic estimate $G(\theta_t)$ for $\nabla_{\theta} J(\theta_t)$, where $G(\theta_t)$ is close to the direction of steepest descent $\nabla_{\theta} J(\theta_t)$. The online algorithm in the main theorem does exactly this using $G(\theta_t) = 2(f(\bar{X}_t) - \beta) \nabla f(X_t) \tilde{X}_t$ as a stochastic estimate for $\nabla_{\theta} J(\theta_t)$.

For large t , we expect that $\mathbf{E}[f(\bar{X}_t) - \beta] \approx \mathbf{E}_{Y \sim \pi_{\theta_t}}[f(Y) - \beta]$ and $\mathbf{E}[\nabla f(X_t) \tilde{X}_t] \approx \nabla_{\theta}(\mathbf{E}_{Y \sim \pi_{\theta_t}}[f(Y) - \beta])$ since θ_t is changing very slowly, this is because t becomes large due to $\lim_{t \rightarrow \infty} \alpha_t = 0$. We need to mind that even though the expectations above is fulfilled, in order to make the product of $\mathbf{E}[f(\bar{X}_t) - \beta]$ and $\mathbf{E}[\nabla f(X_t) \tilde{X}_t]$ be the estimate of $\nabla_{\theta} J(\theta_t)$, it is important to highlight that for random variables X and Y , it is not always true that $\mathbf{E}[XY] = \mathbf{E}X \cdot \mathbf{E}Y$ unless X and Y are independent. This is the reason why the process \tilde{X}_t is needed. Since \tilde{X}_t and X_t are driven by independent Brownian motions, we expect that $\mathbf{E}[2(f(\bar{X}_t) - \beta) \nabla f(X_t) \tilde{X}_t] \approx \nabla_{\theta} J(\theta_t)$ for large t due to \bar{X}_t and (X_t, \tilde{X}_t) becoming asymptotically independent since θ_t will be changing very slowly for large t .

Thus, we expect that for large t , the stochastic sample $2(f(\bar{X}_t) - \beta) \nabla f(X_t) \tilde{X}_t$ will provide an asymptotically unbiased estimate for the direction of steepest descent $\nabla_\theta J(\theta_t)$ and $\|\nabla_\theta J(\theta_t)\|$ will converge to zero as $t \rightarrow \infty$. The above analysis is just some intuitive inspirations, what we need is to make sure this works. Hence rigorous mathematical proofs in the original paper is needed.

2.3 Summary of the Mathematical Proof

The original paper [4] focuses on proving the convergence properties of an online algorithm for optimizing over the stationary distribution of ergodic stochastic differential equations (SDEs). Key mathematical results are developed to ensure the algorithm's efficacy, including the analysis of Poisson partial differential equations (PDEs), bounds on the moments of the stochastic processes involved, and the demonstration of convergence to a stationary point.

2.3.1 Proposition

Statement: Vital proposition in the original paper provides a bound for the convergence of the proposed algorithm under certain regularity conditions.

Proof Outline:

1. **Lemma A.1:** Establishes a bound for the difference of exponential terms involving the process states.

- For $m', k \in \mathbb{R}^+$, constants $C, m > 0$ exist such that for $x, x' \in \mathbb{R}^d$:

$$e^{-|x'-x|^2} \cdot |x' - x|^k \leq \frac{C}{(1 + |x|^m)(1 + |x'|^{m'})}$$

2. **Convergence Analysis:**

- Utilizes the closed-form solution for the distribution of the process.
- Defines functions $f(t, x, \theta)$ and $\Sigma_t(\theta)$ to represent the mean and covariance of the process X_t^θ .
- Shows that the stationary distribution $X_t^\theta \sim N(f(t, x, \theta), \Sigma_t(\theta))$ leads to:

$$X_t^\theta \sim N(h^{-1}(\theta)g(\theta), \sigma^2(2h(\theta))^{-1})$$

- By leveraging properties of positive definite matrices and the eigenvalues of $\Sigma_t(\theta)$, the proof demonstrates the boundedness and convergence of the algorithm.

2.3.2 Theorem 2.2 of the original paper [4]

Statement: Theorem 2.2 asserts that under Assumption 2.1, the algorithm converges to a stationary point almost surely.

Proof Outline:

1. Decomposition:

- Decomposes the evolution of θ_t into a direction of steepest descent and two fluctuation terms.
- Defines error terms Z_t^1 and Z_t^2 representing fluctuations due to the stochastic nature of the process.

2. Stopping Times:

- Introduces cycles of stopping times to control periods when the gradient $|\nabla_{\theta} J(\theta_t)|$ is close to zero and when it is not.
- Uses Lemmas 3.9 and 3.10 to show that for sufficiently large k , the differences $J(\theta_{\sigma_k}) - J(\theta_{\tau_k})$ and $J(\theta_{\tau_k}) - J(\theta_{\sigma_{k-1}})$ are bounded by constants γ_1 and γ_2 , respectively, ensuring convergence.

3. Proof by Contradiction:

- Assumes an infinite number of stopping times τ_k and shows that the cumulative decrease in $J(\theta_t)$ contradicts the boundedness of the function, thereby proving convergence.

The rigorous mathematical analysis provides a solid foundation for the convergence claims, combining probabilistic methods, PDE analysis, and stochastic process theory to establish the effectiveness of the proposed optimization algorithm.

2.4 Computational Limitations

There are several theoretical limitations of the proposed algorithm:

2.4.1 Dependence on Hyper-parameters

The performance of the algorithm is sensitive to the selection of hyper-parameters such as the learning rate and mini-batch size. The learning rate must decay as $t \rightarrow \infty$, but it should not decrease too rapidly, and the initial magnitude should be sufficiently large to ensure convergence.

2.4.2 Stochastic Estimates and Noise

The gradient descent direction's estimation is stochastic and depends on the mini-batch size. A larger mini-batch size reduces the noise in the estimation but may increase computational cost. The algorithm's effectiveness in reducing this noise is crucial for its performance.

2.4.3 Convergence to Global Minima

The algorithm may converge to a local minimizer. This is true of all optimization algorithms, not just the online forward propagation algorithm. For problems with a unique global minimizer, the algorithm converges to the optimal solution if the learning rate is chosen correctly. However, for problems with multiple global minimizers, the algorithm may converge to any one of the global minimizers, depending on the initial conditions and the learning rate schedule.

2.4.4 High-dimensional Problems

While the algorithm is designed to handle high-dimensional stochastic optimal control problems, it may face computational challenges in very high dimensions. Traditional numerical methods struggle with high-dimensional PDEs, and although the proposed algorithm aims to address this, its scalability and efficiency in extremely high dimensions remain a concern.

2.4.5 Path-dependent and Nonlinear SDEs

The algorithm's implementation and performance vary for different types of SDE models, including path-dependent and nonlinear SDEs. The theoretical guarantees provided for linear SDEs may not fully extend to these more complex models.

2.4.6 Application-specific Challenges

The algorithm is tested on applications in mathematical finance, such as parameter estimation for SDE models and stochastic optimal control. Some assumptions may not hold in all practical scenarios, potentially affecting the algorithm's applicability and performance. Each application presents unique challenges, and the algorithm's performance may vary based on the specific characteristics and requirements of each application.

Chapter 3

Implementation

3.1 Simple Cases

3.1.1 One-Dimensional Ornstein-Uhlenbeck Process To Start With

$$dX_t^\theta = (\theta - X_t^\theta) dt + dW_t \quad (3.1)$$

In this case, the main algorithm is used to learn the minimizer for $J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y - 2)^2$. Note that in this case we have the closed-form solution $\pi_\theta \sim N(\theta, \frac{1}{2})$ and thus the global minimizer is $\theta^* = 2$.

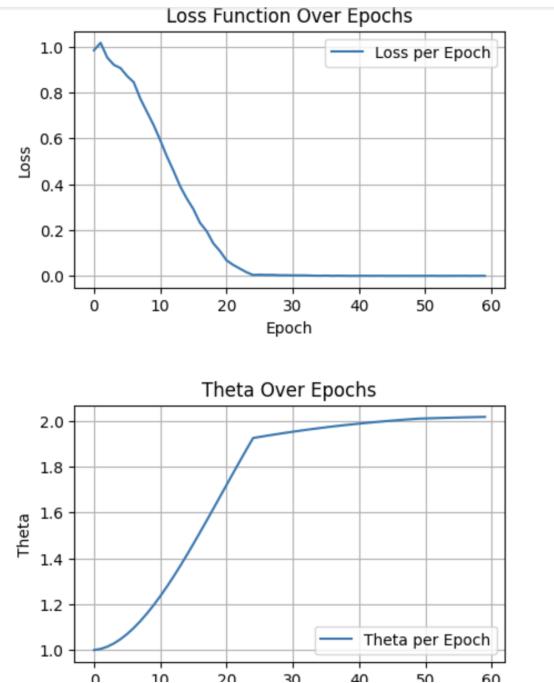


Figure 3.1: $J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y - 2)^2$

The result of the implementation is above. With total epochs equal to 60, final theta is $\theta = 2.016258716583252$ and final loss $J(\theta) = 0.0002$, which is consistent with the theory. To show the algorithm indeed converge in the long run. We adjust the total number of epochs to be 100. The below results shows that the algorithm indeed converge.

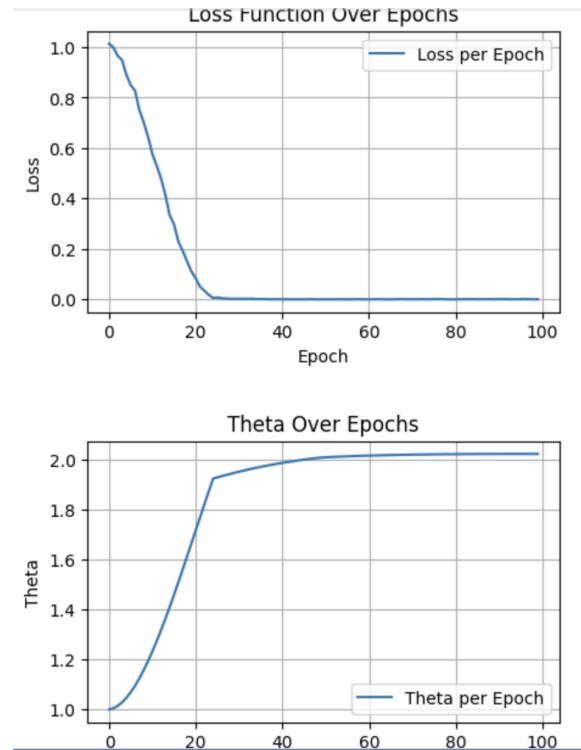


Figure 3.2: $\#epochs = 100$

Similarly, we also use the algorithm to learn the minimizer for $J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y^2 - 2)^2$. In this case, the two global minimizers are $\theta^* = \pm\sqrt{1.5}$.

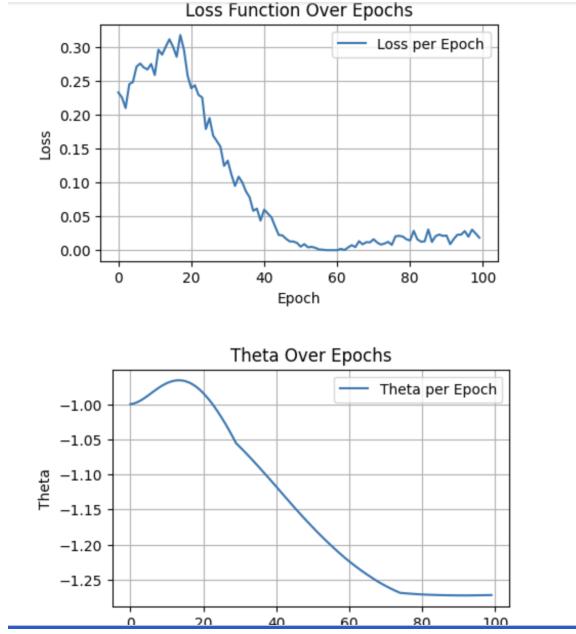


Figure 3.3: $J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y^2 - 2)^2$

Here the parameter trained by the online algorithm converges to a global minimizer. The global minimizer which the algorithm converges to depends on the initial value of θ .

More General OU Processes

$$dX_t^\theta = (\theta^1 - \theta^2 X_t^\theta) dt + dW_t \quad (3.2)$$

In this case, $\theta = (\theta^1, \theta^2)$.

The objective function $J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y^2 - 2)^2$. Algorithm with batch size N is as follows:

$$\begin{aligned} d\theta_t^1 &= -4\alpha_t \left(\frac{1}{N} \sum_{i=1}^N \left(\bar{X}_t^{(i)} \right)^2 - 2 \right) \cdot \left(\frac{1}{N} \sum_{i=1}^N X_t^{(i)} \tilde{X}_t^{1,(i)} \right) dt \\ d\theta_t^2 &= -4\alpha_t \left(\frac{1}{N} \sum_{i=1}^N \left(\bar{X}_t^{(i)} \right)^2 - 2 \right) \cdot \left(\frac{1}{N} \sum_{i=1}^N X_t^{(i)} \tilde{X}_t^{2,(i)} \right) dt \\ dX_t^{(i)} &= \left(\theta_t^1 - \theta_t^2 X_t^{(i)} \right) dt + dW_t^i \\ d\tilde{X}_t^{1,(i)} &= \left(1 - \theta_t^2 \tilde{X}_t^{1,(i)} \right) dt \\ d\tilde{X}_t^{2,(i)} &= \left(-X_t^{(i)} - \theta_t^2 \tilde{X}_t^{2,(i)} \right) dt \\ d\bar{X}_t^{(i)} &= \left(\theta_t^1 - \theta_t^2 \bar{X}_t^{(i)} \right) dt + d\bar{W}_t^i \end{aligned} \quad (3.3)$$

for $i = 1, 2, \dots, N$.

In practice, we choose the batch size $N = 10000$, so the training can be more stable. The below figure shows the result of this two parameter case.

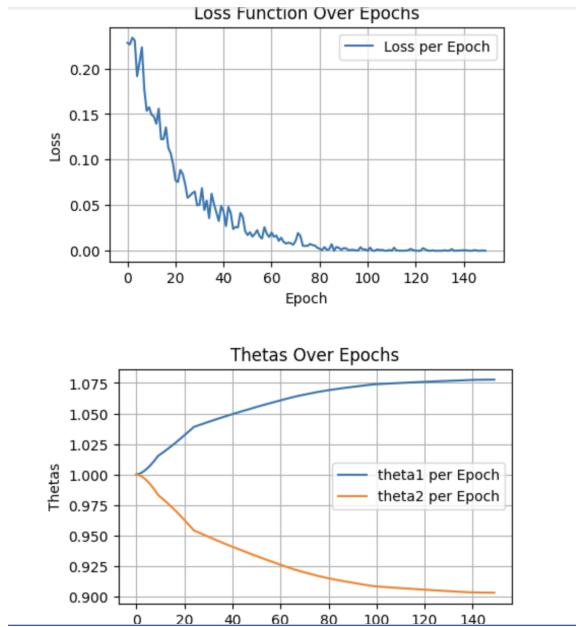


Figure 3.4: Two Parameter case

Final theta1 is tensor(1.0776), final theta2 equals to tensor(0.9033), and final $J(\theta)$ is $4.1243e - 06$ which is small.

NOTE: Computation of $E_{\pi_\theta}[Y^2]$

To compute $E_{\pi_\theta}[Y^2]$, it is obvious we have to natural methods. The first method is as follows.

```
import torch
import matplotlib.pyplot as plt

# Simulate the OU process of e.g. 4.1 two parameter case
# X_simu = torch.randn(N, requires_grad=False)
batch_size = 10000
X_simu = torch.ones(batch_size)
dt_simu = 0.1
number_of_step = 10000
# theta1 = torch.randn(1)
```

```

# theta2 = torch.randn(1)
theta1 = 1.075
theta2 = 0.9

for i in range(1, number_of_step):
    dW = torch.randn(batch_size) * torch.sqrt(torch.tensor(dt_simu))
    X_simu = X_simu + (theta1 - theta2 * X_simu) * dt_simu + dW

print(torch.mean(X_simu**2))

```

We have another method. This is using the ergodic theorem.

$$\lim_{t \rightarrow \infty} \frac{1}{t} \int_0^t f(X_s^\theta) ds = \mathbf{E}_{Y \sim \pi_\theta} f(Y) \quad \text{a.s. ,}$$

and the implementation is below.

```

import torch

# Simulate the OU process
batch_size = 10000
X_simu = torch.ones(batch_size)
dt_simu = 0.01
number_of_step = 100000
# theta1 = torch.randn(1)
# theta2 = torch.randn(1)
theta1 = 1.075
theta2 = 0.9
estimator = 0

for i in range(1, number_of_step):
    dW = torch.randn(batch_size) * torch.sqrt(torch.tensor(dt_simu))
    delta_X_simu = (theta1 - theta2 * X_simu) * dt_simu + dW
    X_simu = X_simu + delta_X_simu
    estimator += (1/(number_of_step * dt_simu)) * X_simu**2 * dt_simu

print(torch.mean(estimator))

```

We conclude that both methods have the same result. The only thing to consider is the so called cost of implementation. After consideration we choose to use the first method since it is cheaper.

3.1.2 One-Dimensional Nonlinear Process

We now try to use the online algorithm to optimize over the stationary distribution of a one-dimensional nonlinear process

$$dX_t^\theta = \left(\theta - X_t^\theta - (X_t^\theta)^3 \right) dt + dW_t \quad (3.4)$$

Note that the format of the drift term is not the same as what is required in the main theorem. But we can still use the main algorithm to learn the minimizer of $J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y^2 - 2)^2$. The following mini-batch algorithm is used:

$$\begin{aligned} d\theta_t &= -4\alpha_t \left(\frac{1}{N} \sum_{i=1}^N \left(\bar{X}_t^{(i)} \right)^2 - 2 \right) \cdot \left(\frac{1}{N} \sum_{i=1}^N X_t^{(i)} \tilde{X}_t^{(i)} \right) dt \\ dX_t^{(i)} &= \left(\theta_t - X_t^{(i)} - (X_t^{(i)})^3 \right) dt + dW_t^{(i)} \\ d\tilde{X}_t^{(i)} &= \left(1 - \tilde{X}_t^{(i)} - 3(X_t^{(i)})^2 \tilde{X}_t^{(i)} \right) dt \\ d\bar{X}_t^{(i)} &= \left(\theta_t - \bar{X}_t^{(i)} - (\bar{X}_t^{(i)})^3 \right) dt + d\bar{W}_t^{(i)} \end{aligned}$$

for $i = 1, 2, \dots, N$.

Figure 3.5 below shows the convergence of the parameter θ_t . And it shows the loss function decays to zero, i.e. the global minimum, very quickly. The final theta is 4.351974964141846, and final J is tensor(0.0014, device='cuda:0'), which is in a acceptable range. Hence we see the interesting thing: Even though the requirement is not totally fulfilled in some problems, it is possible that our algorithm can be adapted to those problem.

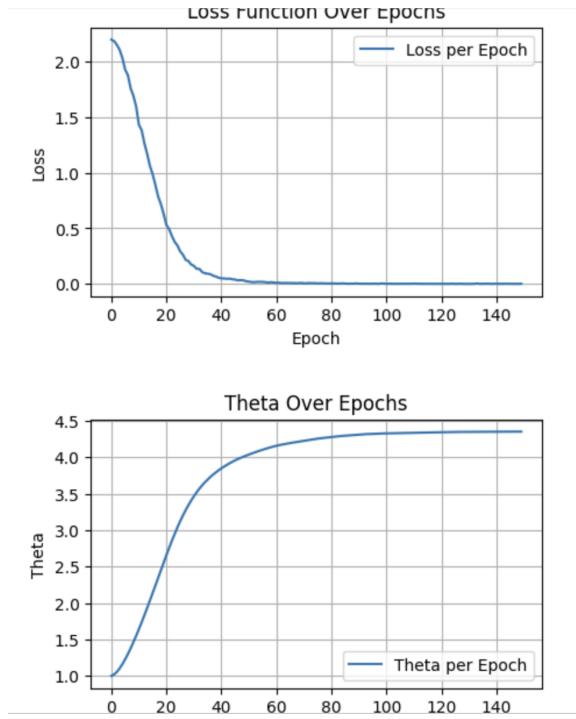


Figure 3.5: Non linear case

3.1.3 Optimizing over the Drift and Volatility Coefficients

Linear Case

We now optimize over the drift and volatility functions of the process

$$dX_t^\theta = (\mu - X_t^\theta) dt + \sigma dW_t \quad (4.11)$$

with parameters $\theta = (\mu, \sigma)$.

Since the format perfectly satisfies the requirement, our online algorithm can be used to learn the minimizer of $J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y^2 - 2)^2$. The variant is written below:

$$\begin{aligned}
d\mu_t &= -4\alpha_t \left(\frac{1}{N} \sum_{i=1}^N \left(\bar{X}_t^{(i)} \right)^2 - 2 \right) \cdot \left(\frac{1}{N} \sum_{i=1}^N X_t^{(i)} \tilde{X}_t^{1,(i)} \right) dt \\
d\sigma_t &= -4\alpha_t \left(\frac{1}{N} \sum_{i=1}^N \left(\bar{X}_t^{(i)} \right)^2 - 2 \right) \cdot \left(\frac{1}{N} \sum_{i=1}^N X_t^{(i)} \tilde{X}_t^{2,(i)} \right) dt \\
dX_t^i &= \left(\mu_t - X_t^{(i)} \right) dt + \sigma_t dW_t^{(i)} \\
d\tilde{X}_t^{1,(i)} &= \left(1 - \tilde{X}_t^{1,(i)} \right) dt \\
d\tilde{X}_t^{2,(i)} &= -\tilde{X}_t^{2,(i)} dt + dW_t^{(i)} \\
d\bar{X}_t^{(i)} &= \left(\mu_t - \bar{X}_t^{(i)} \right) dt + \sigma_t d\bar{W}_t^{(i)}
\end{aligned}$$

for $i = 1, 2, \dots, N$.

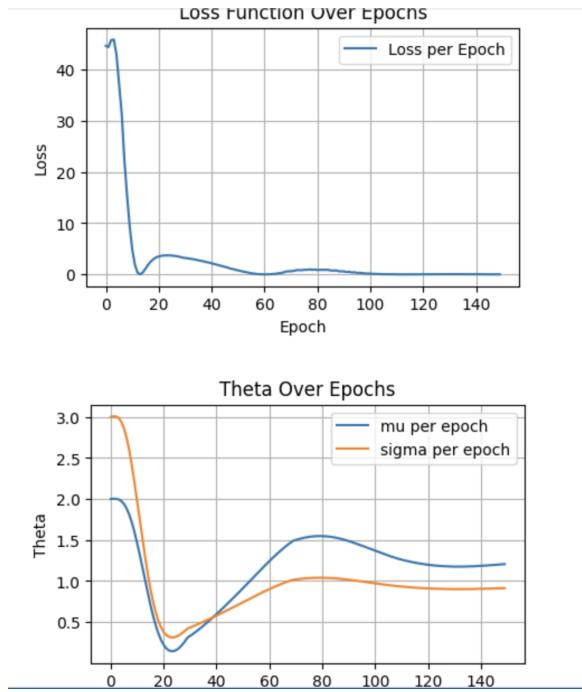


Figure 3.6: Drift and Volatility Linear

In Figure 3.6, the trained parameters μ_t, σ_t converge and the objective function $J(\theta_t) \rightarrow 0$ very quickly. And we have results Final mu: 1.2057600021362305; Final sigma: 0.9127710461616516; Final J: 0.011372823268175125.

Notice that there weird things happened. We know $\mu = 1.0$ and $\sigma = 1.5$ is also a set of parameters which can make loss $J(\theta)$ small. This is shown by the output below:

Check Simulation

```

mu = 1.0
sigma = 1.5
# Simulate the OU process
batch_size_simu = 10000
dt_simu = 0.01
number_of_stop_simu = 10000
X_simu = torch.FloatTensor(batch_size_simu, device=device)
loss = 0.0
for i in range(0, number_of_stop_simu):
    X_simu = torch.FloatTensor(batch_size_simu, device=device) + torch.sqrt(torch.tensor(dt_simu, device=device))
    for l in range(0, number_of_stop_simu):
        dt = torch.FloatTensor(batch_size_simu, device=device) * torch.sqrt(dt_simu, device=device)
        X_simu += (mu - X_simu) * dt_simu + sigma * dt
    loss += (torch.mean(X_simu**2) - 2) ** 2
print(loss)

```

Python

Figure 3.7: Drift and Volatility Linear

Hence in this special case there are multiple local minima.

Nonlinear Case

We also implement the online algorithm for the nonlinear process to see the ability to adapt.

$$dX_t^\theta = \left(\mu - (X_t^\theta)^3 \right) dt + \sigma X_t^\theta dW_t \quad (3.5)$$

where $\theta = (\mu, \sigma)$ are the parameters and the objective function is $J(\theta) = (\mathbf{E}_{Y \sim \pi_\theta} Y^2 - 10)^2$. The mini-batch algorithm 3.3 now becomes:

$$\begin{aligned}
d\mu_t &= -4\alpha_t \left(\frac{1}{N} \sum_{i=1}^N \left(\bar{X}_t^{(i)} \right)^2 - 2 \right) \cdot \left(\frac{1}{N} \sum_{i=1}^N X_t^{(i)} \tilde{X}_t^{1,(i)} \right) dt \\
d\sigma_t &= -4\alpha_t \left(\frac{1}{N} \sum_{i=1}^N \left(\bar{X}_t^{(i)} \right)^2 - 2 \right) \cdot \left(\frac{1}{N} \sum_{i=1}^N X_t^{(i)} \tilde{X}_t^{2,(i)} \right) dt \\
dX_t^i &= \left(\mu_t - \left(X_t^{(i)} \right)^3 \right) dt + \sigma_t X_t^{(i)} dW_t^{(i)} \\
d\tilde{X}_t^{1,(i)} &= \left(1 - 3 \left(X_t^{(i)} \right)^2 \tilde{X}_t^{1,(i)} \right) dt + \sigma_t \tilde{X}_t^{1,(i)} dW_t^{(i)} \\
d\tilde{X}_t^{2,(i)} &= -3 \left(X_t^{(i)} \right)^2 \tilde{X}_t^{2,(i)} dt + \left(X_t^{(i)} + \sigma_t \tilde{X}_t^{2,(i)} \right) dW_t^{(i)} \\
d\bar{X}_t^{(i)} &= \left(\mu_t - \left(\bar{X}_t^{(i)} \right)^3 \right) dt + \sigma_t \bar{X}_t^{(i)} d\bar{W}_t^{(i)}
\end{aligned}$$

for $i = 1, 2, \dots, N$.

By calibration (use the final parameters to simulate process and check the loss) we can find the original paper has some typos in this part. And the results I get are described further below.

In Figure 3.8, the trained parameters μ_t, σ_t converge and it also shows that the objective function $J(\theta_t) \rightarrow 0$ very quickly. Hence we proved that the algorithm can adapt this case.

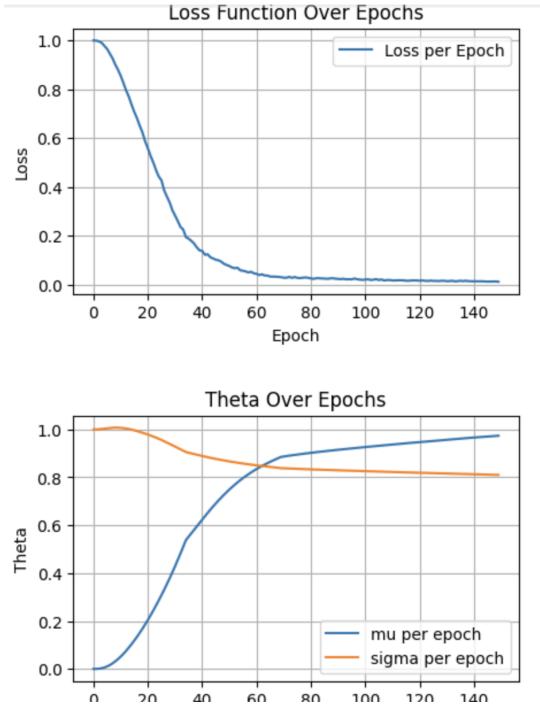


Figure 3.8: Drift and Volatility Non Linear

3.1.4 Multi-Dimensional Ornstein-Uhlenbeck Process

Independent OU Process

We next consider a simple multi-dimensional Ornstein-Uhlenbeck process, it consists of m independent copies of an OU process. For the parameter $\theta = (\theta^1, \theta^2) \in R^{2m}$, let the m-dimensional Ornstein-Uhlenbeck process be

$$dX_t^\theta = (\theta^1 - \theta^2 \odot X_t^\theta) dt + dW_t \quad (3.6)$$

where $X_t^\theta \in R^m$, $W_t \in R^m$, and \odot is an element-wise product.

The objective function is

$$J(\theta) := \left(\sum_{k=1}^m \mathbf{E}_{Y \sim \pi_\theta} |Y_k|^2 - 2m \right)^2 \quad (3.7)$$

The online algorithm becomes:

$$\begin{aligned}
d\theta_t^1 &= -4\alpha_t \left(|\bar{X}_t|^2 - 2 \right) X_t \odot \tilde{X}_t^1 dt \\
d\theta_t^2 &= -4\alpha_t \left(|\bar{X}_t|^2 - 2 \right) X_t \odot \tilde{X}_t^2 dt \\
dX_t &= (\theta_t^1 - \theta_t^2 \odot X_t) dt + dW_t^i \\
d\tilde{X}_t^1 &= \left(1 - \theta_t^2 \odot \tilde{X}_t^1 \right) dt \\
d\tilde{X}_t^2 &= \left(-X_t - \theta_t^2 \odot \tilde{X}_t^2 \right) dt \\
d\bar{X}_t &= (\theta_t^1 - \theta_t^2 \odot \bar{X}_t) dt + d\bar{W}_t^i
\end{aligned}$$

We implement the algorithm for $m = 3$ and $m = 10$. In Figures 3.9 and 3.10 , the objective functions $J(\theta_t) \rightarrow 0$ as t becomes large.

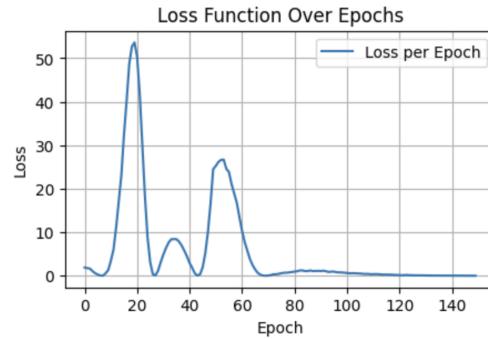


Figure 3.9: Independent $m=3$

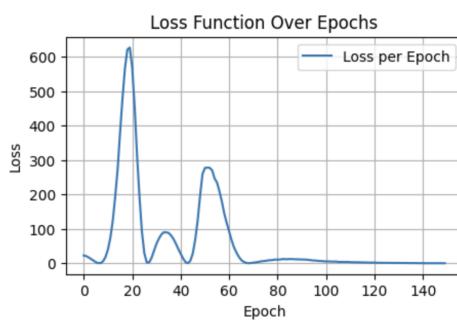


Figure 3.10: Independent $m=10$

It is worth mentioning that the final thetas values and the near-zero final loss suggest effective convergence. And despite the higher initial values, the loss converges

effectively, reaching near-zero levels after epoch 80, showcasing the robustness of the optimization process even with increased dimensionality. Overall, the gradient descent algorithm is working effectively, successfully minimizing the loss and converging to optimal parameter values despite initial instabilities.

Correlated OU Process

For the parameters $\theta = (\mu, \sigma)$ with $\mu \in \mathbb{R}^m, \sigma \in \mathbb{R}^{m \times m}$, let the m -dimensional process X_t^θ satisfy

$$dX_t^\theta = (\mu - X_t^\theta) dt + \sigma dW_t$$

where $W_t \in \mathbb{R}^m$.

Let $X_t^{\theta,i}$ denote the i -th element of X_t^θ and define \tilde{X}_t^μ and \tilde{X}_t^σ as the Jacobian matrices of X_t^θ with respect to μ and σ :

$$\begin{aligned}\tilde{X}_t^\mu &= \nabla_\mu X_t^\theta \in \mathbb{R}^{m \times m}, & \tilde{X}_t^{\mu,i} &= \nabla_\mu X_t^{\theta,i} \in \mathbb{R}^m \\ \tilde{X}_t^\sigma &= \nabla_\sigma X_t^\theta \in \mathbb{R}^{m \times m \times m}, & \tilde{X}_t^{\sigma,i} &= \nabla_\sigma X_t^{\theta,i} \in \mathbb{R}^{m \times m}\end{aligned}$$

for $i \in \{1, 2, \dots, m\}$.

$$dX_t^{\theta,i} = (\mu_i - X_t^{\theta,i}) dt + \sum_j \sigma_{i,j} dW_t^j$$

Now in this case the algorithm 2.3 becomes

$$\begin{aligned}d\mu_t &= -4\alpha_t \left(|\bar{X}_t|^2 - 2m \right) \left(\sum_{k=1}^m X_t^k \tilde{X}_t^{\mu,k} \right) dt \\ d\lambda_t &= -4\alpha_t \left(|\bar{X}_t|^2 - 2m \right) \left(\sum_{k=1}^m X_t^k \tilde{X}_t^{\lambda,k} \right) dt \\ dX_t &= (\mu_t - X_t) dt + \sigma_t dW_t \\ d\bar{X}_t &= (\mu_t - \bar{X}_t) dt + \sigma_t d\bar{W}_t \\ d\tilde{X}_t^\mu &= \left(I_m - \tilde{X}_t^\mu \right) dt \\ d\tilde{X}_t^{\sigma,i} &= -\tilde{X}_t^{\sigma,i} dt + D_i(dW_t), \quad i \in \{1, \dots, m\}\end{aligned}$$

where I_m is the $m \times m$ identity matrix and where $D_i(dW_t)$ is a $m \times m$ matrix with all elements equal to 0 except i -th column being dW_t . We examine the algorithm's performance for dimensions $m = 3, 10$.

In Figures 3.11, the objective function $J(\theta_t)$ does not converge to zero, which indicates there may be issues in the algorithm or the choice of hyper-parameters.

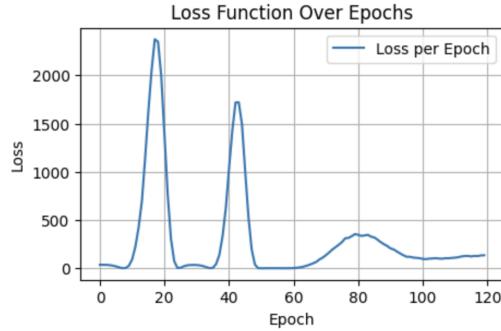
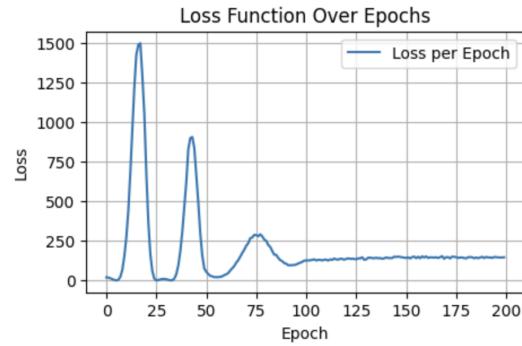


Figure 3.11: Correlated m=3

While the algorithm shows promising signs of effectiveness, addressing the identified instabilities and refining the training process will be essential to achieving more reliable and optimal performance. Hence we refined the code by tuning the learning rate. Unfortunately the loss did not converge to 0 even after the refine.



```
Final mu: tensor([0.7447, 0.7694, 0.7118], device='cuda:0')
Final sigma: tensor([[3.2132, 0.0000, 0.0000],
                      [0.0000, 3.2349, 0.0000],
                      [0.0000, 0.0000, 3.2303]], device='cuda:0')
Final J: tensor(146.1420, device='cuda:0')
```

Figure 3.12: Correlated m=3 refined

And for higher dimensions, $m = 10$, the algorithm loss dose not converge to 0.

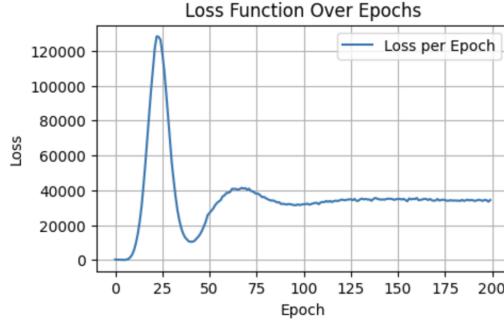


Figure 3.13: Correlated m=10 refined

Notice that our mathematical proof only ensures that $\lim_{t \rightarrow \infty} |\nabla_{\theta} J(\theta_t)| \stackrel{\text{a.s.}}{\equiv} 0$, while it dose not guaranty the loss to converge to 0.

3.1.5 Path-dependent SDE

We consider the path-dependent SDE

$$dX_t^{\theta} = \left(\theta - X_t^{\theta} - \frac{1}{t} \int_0^t X_s^{\theta} ds \right) dt + dW_t \quad (3.8)$$

where $X_t^{\theta}, W_t \in \mathbb{R}$.

Although path-dependent SDEs are not directly addressed by this article's convergence theory, the numerical example in this subsection suggests that the online forward propagation algorithm can also be applied to path-dependent stochastic processes.

For this numerical example, the objective function is $J(\theta) = (\mathbf{E}_{Y \sim \pi_{\theta}} Y - 2)^2$. The SDE 3.8 does not fit the problem described in 2.1. However, our algorithm still can find the global optimum. Which shows the ability of adapt. Now the online algorithm 2.3 is:

$$\begin{aligned} d\theta_t &= -4\alpha_t (\bar{X}_t - 2) \tilde{X}_t dt \\ dX_t &= \left(\theta_t - X_t - \frac{1}{t} \int_0^t X_s ds \right) dt + dW_t \\ d\tilde{X}_t &= \left(1 - \tilde{X}_t - \frac{1}{t} \int_0^t \tilde{X}_s ds \right) dt \\ d\bar{X}_t &= \left(\theta_t - \bar{X}_t - \frac{1}{t} \int_0^t \bar{X}_s ds \right) dt + d\bar{W}_t \end{aligned}$$

In Figure 3.14 , the trained parameter converges. The objective function $J(\theta_t)$ is approximated using a time-average. And the objective function $J(\theta_t)$ converges to 0 very quickly.

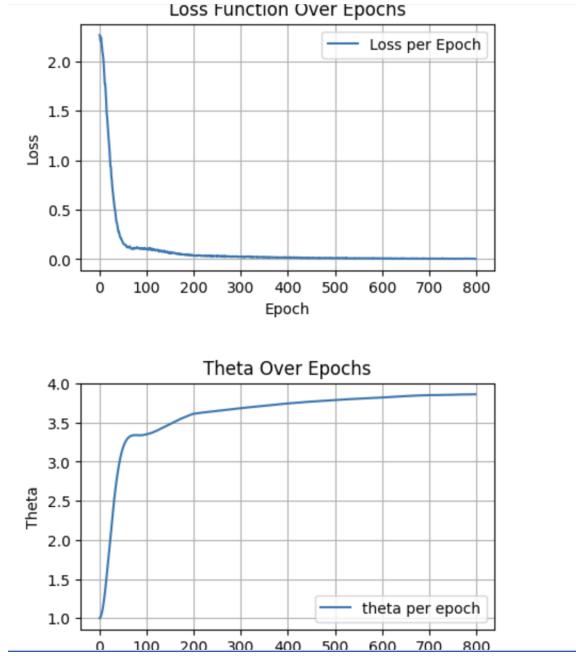


Figure 3.14: Path Dependent

3.2 Linear–quadratic Regulator

The online optimization algorithm can be employed to solve stochastic optimal control problems, including high-dimensional scenarios where traditional numerical methods, such as solving the Hamilton-Jacobi-Bellman (HJB) equation with finite difference methods, are computationally expensive or impractical. As a numerical example, we consider the classic Linear Quadratic Regulator (LQR) problem, which has numerous financial applications such as optimal execution[1]. Let $\{X_t\}_{t \geq 0}$ be the state process governed by the stochastic differential equation (SDE):

$$dX_t = (AX_t + BU_t) dt + \sigma dW_t \quad (3.9)$$

where $X_0 = x_0$, $X_t \in \mathbb{R}^n$, and matrices $A, \sigma \in \mathbb{R}^{n \times n}$, $B \in \mathbb{R}^{n \times m}$.

Here, $\{W_t\}_{t \geq 0}$ is an \mathbb{R}^n -valued standard Wiener process, and $\{U_t\}_{t \geq 0} \in \mathbb{R}^m$ represents the control. The objective is to learn a control process U_t to minimize the ergodic cost functional for system below:

$$J(U) = \lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T (X_t^T Q X_t + U_t^T R U_t) dt \quad (3.10)$$

where Q and R are positive definite matrices. It is well-known that the optimal control is given by: $U = -R^{-1}B^\top K X$, where K is the unique solution of the following algebraic Riccati equation (ARE):

$$A^\top K + KA - KBR^{-1}B^\top K + Q = 0 \quad (3.11)$$

To evaluate the accuracy of our algorithm in solving stochastic optimal control problems, we applied it to several high-dimensional stochastic LQR problems. The LQR problem was chosen due to its closed-form solution, even in high dimensions, which provides a reliable benchmark for assessing our algorithm's accuracy. We present a series of numerical examples demonstrating how the online optimization algorithm learns parametric controls for various LQR problems. These parametric controls are implemented as either linear functions or neural networks.

3.2.1 One-dimensional Linear Control

For the first step, we implement the online optimization algorithm for the one-dimensional case with a linear control function. For simplicity, we assume that $A = -1, B = \sigma = Q = R = 1$ for 3.9, then we get:

$$\begin{aligned} dX_t^\theta &= (-X_t^\theta + \theta X_t^\theta) dt + dW_t \\ J(\theta) &= \lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T (1 + \theta^2) (X_t^\theta)^2 dt \end{aligned}$$

The system becomes

$$\begin{aligned} d\theta_t &= -\alpha_t \left[\frac{1}{N} \sum_{i=1}^N \left(2\theta_t \left(X_t^{(i)} \right)^2 + 2(1 + \theta_t^2) X_t^{(i)} \tilde{X}_t^{(i)} \right) \right] dt \\ dX_t^{(i)} &= (\theta_t - 1) X_t^{(i)} dt + dW_t^{(i)} \\ d\tilde{X}_t^{(i)} &= \left(X_t^{(i)} + (\theta_t - 1) \tilde{X}_t^{(i)} \right) dt \end{aligned}$$

with $i = 1, 2, \dots, N$.

Solving the ARE yields the optimal control $\theta^* = -0.41421$. Here we can directly use loss to measure the performance. Figure 3.15 below shows that the parameter θ_t trained with the online optimization algorithm converges to θ^* .

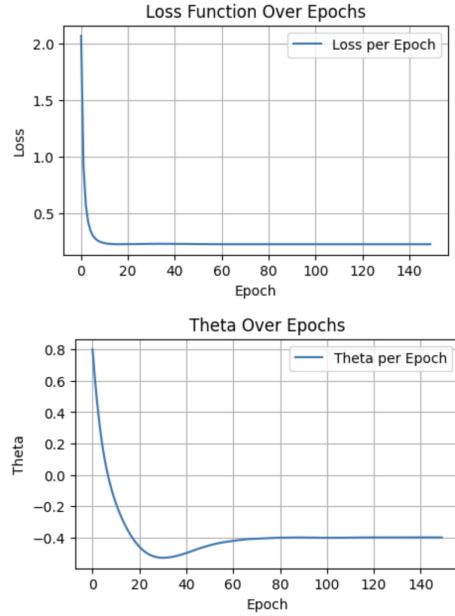


Figure 3.15: LQR 1-D

The code is given in appendix A.3

3.2.2 Multi-dimensional Linear Control

We next solve a multi-dimensional LQR problem with a linear control function. For simplicity, we assume that $m = n, A = -I_n, B = \sigma = I_n$ in 3.9 where I_n is n dimensional identity matrix. That is,

$$dX_t^\theta = (-X_t^\theta + \theta X_t^\theta) dt + dW_t \quad (3.12)$$

$$J(\theta) = \lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T (X_t^\theta)^\top (Q + \theta^\top R \theta) X_t^\theta dt \quad (3.13)$$

where $\theta \in \mathbb{R}^{n \times n}$.

Let $X_t^{\theta,i}$ denote the i -th element of X_t^θ and define

$$\tilde{X}_t^\theta = \nabla_\theta X_t^\theta, \quad \tilde{X}_t^{\theta,i} = \nabla_\theta X_t^{\theta,i}, \forall i \in \{1, 2, \dots, n\}$$

\tilde{X}_t^θ has dimensions $n \times n \times n$ and $\tilde{X}_t^{\theta,i}$ has dimensions $n \times n$. Note that when we are training over a mini-batch of size N , \tilde{X}_t^θ has dimensions $N \times n \times n \times n$.

We first discuss the methods necessary for the computationally efficient simulation of the gradient $\nabla_\theta X_t^\theta$. The state process from 3.12 satisfies

$$dX_t^{\theta,i} = \left(-X_t^{\theta,i} + \sum_{j=1}^n \theta_{i,j} X_t^{\theta,j} \right) dt + dW_t^i$$

and therefore

$$d\tilde{X}_t^{\theta,i} = \left(-\tilde{X}_t^{\theta,i} + \sum_{j=1}^n \theta_{i,j} \tilde{X}_t^{\theta,j} + D_i(X_t^\theta) \right) dt$$

where $D_i(X_t^\theta)$ is an $n \times n$ matrix whose elements are all zeros except for the i -th row, which has values X_t^θ . The gradient of the objective function in 3.13 is:

$$\begin{aligned} \nabla_\theta \left[(X_t^\theta)^\top (Q + \theta^\top R \theta) X_t^\theta \right] &= \sum_{i,j} \nabla_\theta (q_{i,j} + \theta_{:,i}^\top R \theta_{:,j}) X_t^{\theta,i} X_t^{\theta,j} \\ &\quad + 2 \sum_{i,j} \nabla_\theta X_t^{\theta,i} (q_{i,j} + \theta_{:,i}^\top R \theta_{:,j}) X_t^{\theta,j} \\ &= \sum_{i,j} ((R\theta)_{k,j} 1_{\{\ell=i\}} + (R\theta)_{k,i} 1_{\{\ell=j\}})_{n \times n} X_t^{\theta,i} X_t^{\theta,j} \\ &\quad + 2 \sum_{i,j} \tilde{X}_t^{\theta,i} (q_{i,j} + \theta_{:,i}^\top R \theta_{:,j}) X_t^{\theta,j}, \end{aligned}$$

where $((R\theta)_{k,j} 1_{\{\ell=i\}} + (R\theta)_{k,i} 1_{\{\ell=j\}})_{n \times n}$ denotes for a $n \times n$ matrix whose k -th row and ℓ -th column is $(R\theta)_{k,j} 1_{\{\ell=i\}} + (R\theta)_{k,i} 1_{\{\ell=j\}}$

In multi-dimensional case, the performance can not be measured by the loss function $J(\theta)$ any more, since the numerical computation of loss J is hard. This is because errors show up while computing the integral $\int_0^T (X_t^\theta)^\top (Q + \theta^\top R \theta) X_t^\theta dt$ using numerical discretization. For a fixed parameter θ , the time-averaged objective function should converge as $t \rightarrow \infty$. "the Riemann Upper Sum" in this case is only close to the true value of the integral when the time step dt is very small.

Hence we use a different way. The error metrics are defined as:

$$\begin{aligned} \text{Ave Error} &= \frac{\sum_{i,j=1}^n |\theta_{t,i,j} - \theta_{i,j}^*|}{\sum_{i,j=1}^n |\theta_{i,j}^*|} \\ \text{Max Error} &= \frac{\max_{i,j \in \{1, 2, \dots, n\}} |\theta_{t,i,j} - \theta_{i,j}^*|}{\frac{1}{n^2} \sum_{i,j=1}^n |\theta_{i,j}^*|} \\ \text{Cost Error} &= \frac{|J(\theta_T) - J(\theta^*)|}{|J(\theta^*)|} \end{aligned}$$

where θ^* is the optimal control and θ_t is the parameter during training. $J(\theta_T)$ and $J(\theta^*)$ denote the objective function $J(\theta)$ in 3.13 with the parameters θ_T and θ^* , respectively.

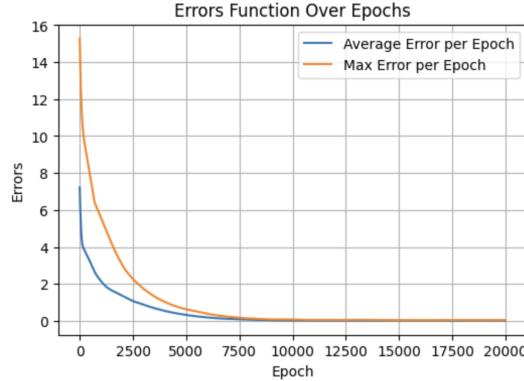


Figure 3.16: LQR Multi-D n=5

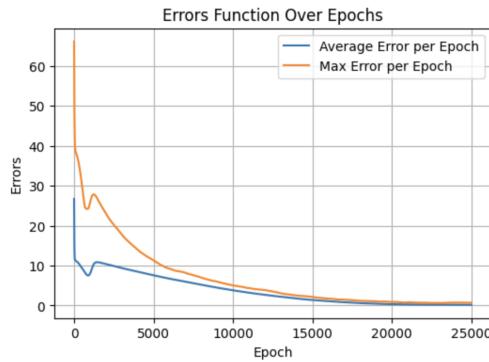


Figure 3.17: LQR Multi-D n=20

3.3 Neural Network Control

3.3.1 One-dimensional Neural Network Control

Now let's train a single-layer neural network control using the online optimization algorithm. The state process is:

$$dX_t^\theta = (-X_t^\theta + f_\theta(X_t^\theta)) dt + dW_t$$

where the control $f_\theta(\cdot)$ is a single-layer neural network $f_\theta(x) = \sum_{i=1}^m c^i \sigma(w^i x + b^i)$, with parameters $\theta = (c^i, w^i, b_i)_{i=1}^m$. The objective function is

$$J(\theta) = \lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T (X_t^\theta)^2 + (f_\theta(X_t^\theta))^2 dt \quad (3.14)$$

Define the gradient of X_t with respect to the parameters as:

$$\tilde{X}_t^w = \nabla_w X_t^\theta \in \mathbb{R}^m, \quad \tilde{X}_t^b = \nabla_b X_t^\theta \in \mathbb{R}^m, \quad \tilde{X}_t^c = \nabla_c X_t^\theta \in \mathbb{R}^m$$

The coupled system 3.3 becomes

$$\begin{aligned} dw_t &= -\alpha_t \left(2X_t \tilde{X}_t^w + 2f_{\theta_t}(X_t) \left(c_t \odot \sigma' (w_t X_t + b_t) X_t + f'_{\theta_t}(X_t) \tilde{X}_t^w \right) \right) dt \\ db_t &= -\alpha_t \left(2X_t \tilde{X}_t^b + 2f_{\theta_t}(X_t) \left(c_t \odot \sigma' (w_t X_t + B_t) + f'_{\theta_t}(X_t) \tilde{X}_t^b \right) \right) dt \\ dc_t &= -\alpha_t \left(2X_t \tilde{X}_t^c + 2f_{\theta_t}(X_t) \left(\sigma (w_t X_t + b_t) + f'_{\theta_t}(X_t) \tilde{X}_t^c \right) \right) dt \\ dX_t &= (-X_t + f_{\theta_t}(X_t)) dt + dW_t \\ d\tilde{X}_t^w &= \left(-\tilde{X}_t^w + c_t \odot \sigma' (w_t X_t^i + b_t) X_t + f'_{\theta_t}(X_t) \tilde{X}_t^w \right) dt \\ d\tilde{X}_t^b &= \left(-\tilde{X}_t^b + c_t \odot \sigma' (w_t X_t^i + b_t) + f'_{\theta_t}(X_t) \tilde{X}_t^b \right) dt \\ d\tilde{X}_t^c &= \left(-\tilde{X}_t^c + \sigma (w_t X_t + b_t) + f'_{\theta_t}(X_t) \tilde{X}_t^c \right) dt \\ d\bar{X}_t &= (-\bar{X}_t + f_{\theta_t}(\bar{X}_t)) dt + d\bar{W}_t \end{aligned}$$

In the original paper we use error metrics to measure the performance of the algorithm. But here in the 1-dimensional case, we can directly use the loss function (3.14) to measure the performance. The training result for 1 dimensional LQR with network control is presented in Figure 3.18. We can observe that $J(\theta)$ is decreasing while final J is 0.008304406888782978, which shows the performance is good. And since the code becomes complicated we also check whether the gradient the algorithm gives is consistent with the gradient the finite difference gives. The result is shown in figure 3.19. It shows that those gradients are consistent.

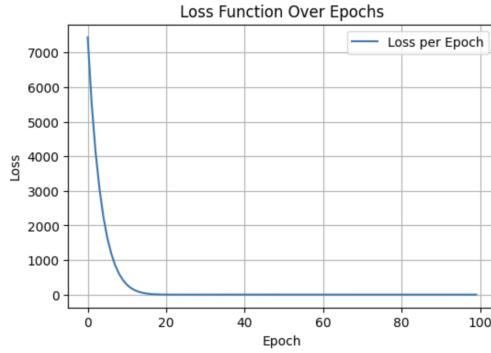


Figure 3.18: NN control 1-D

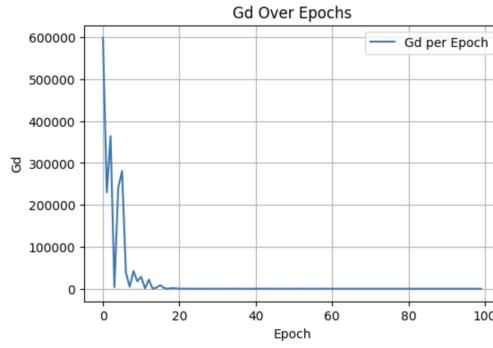


Figure 3.19: NN control 1-D Gradient Check

The code is provided in Appendix A.1.

3.3.2 Multi-dimensional Neural Network Control

We now optimize a single-layer neural network control for a high-dimensional state process:

$$dX_t^\theta = (-X_t^\theta + f_\theta(X_t^\theta)) dt + dW_t \quad (3.15)$$

$$J(\theta) = \lim_{T \rightarrow \infty} \frac{1}{T} \int_0^T (X_t^\theta)^\top Q X_t^\theta + (f_\theta(X_t^\theta))^\top R f_\theta(X_t^\theta) dt \quad (3.16)$$

where $X_t^\theta \in \mathbb{R}^n$ and the single-layer neural network with m hidden units is:

$$f_\theta(x) = c\sigma(wx + b)$$

where $w \in \mathbb{R}^{m \times n}$, $b \in \mathbb{R}^m$, and $c \in \mathbb{R}^{n \times m}$.

As in the below algorithm , define

$$\begin{aligned}\tilde{X}_t^w &= \nabla_w X_t^\theta \in \mathbb{R}^{n \times m \times n}, \quad \tilde{X}_t^{w,i} = \nabla_w X_t^{\theta,i} \in \mathbb{R}^{m \times n} \\ \tilde{X}_t^b &= \nabla_b X_t^\theta \in \mathbb{R}^{n \times m}, \quad \tilde{X}_t^{b,i} = \nabla_b X_t^{\theta,i} \in \mathbb{R}^m \\ \tilde{X}_t^c &= \nabla_c X_t^\theta \in \mathbb{R}^{n \times n \times m}, \quad \tilde{X}_t^{c,i} = \nabla_c X_t^{\theta,i} \in \mathbb{R}^{n \times m}\end{aligned}$$

for $i = 1, 2, \dots, n$.

The online algorithm 2.3 becomes:

$$\begin{aligned}dw_t &= -\alpha_t \left[\nabla_w \left(f_{\theta_t}(X_t)^\top R f_{\theta_t}(X_t) \right) + \sum_{i=1}^n \frac{\partial}{\partial x_i} \left((X_t)^\top Q X_t + f_{\theta_t}(X_t)^\top R f_{\theta_t}(X_t) \right) \tilde{X}_t^{w,i} \right] dt \\ db_t &= -\alpha_t \left[\nabla_b \left(f_{\theta_t}(X_t)^\top R f_{\theta_t}(X_t) \right) + \sum_{i=1}^n \frac{\partial}{\partial x_i} \left((X_t)^\top Q X_t + f_{\theta_t}(X_t)^\top R f_{\theta_t}(X_t) \right) \tilde{X}_t^{b,i} \right] dt \\ dc_t &= -\alpha_t \left[\nabla_c \left(f_{\theta_t}(X_t)^\top R f_{\theta_t}(X_t) \right) + \sum_{i=1}^n \frac{\partial}{\partial x_i} \left((X_t)^\top Q X_t + f_{\theta_t}(X_t)^\top R f_{\theta_t}(X_t) \right) \tilde{X}_t^{c,i} \right] dt \\ dX_t &= (-X_t + f_{\theta_t}(X_t)) dt + dW_t \\ d\tilde{X}_t^{w,i} &= \left(-\tilde{X}_t^{w,i} + \sum_k c_{t,i,k} \sigma' (w_t X_t + b_t)_k \left(\sum_\ell w_{t,k,\ell} \tilde{X}_t^{w,\ell} \right) + (c_{t,i,:})^\top \odot \sigma' (w_t X_t + b_t) (X_t)^\top \right) dt \\ d\tilde{X}_t^{b,i} &= \left(-\tilde{X}_t^{b,i} + \sum_k c_{t,i,k} \sigma' (w_t X_t + b_t)_k \left(\sum_\ell w_{t,k,\ell} \tilde{X}_t^{b,\ell} \right) + (c_{t,i,:})^\top \odot \sigma' (w_t X_t + b_t) \right) dt \\ d\tilde{X}_t^{c,i} &= \left(-\tilde{X}_t^{c,i} + \sum_k c_{t,i,k} \sigma' (w_t X_t + b_t)_k \left(\sum_\ell w_{t,k,\ell} \tilde{X}_t^{c,\ell} \right) + D_i(\sigma(w_t X_t + b_t)) \right) dt \\ d\bar{X}_t &= (-\bar{X}_t + f_{\theta_t}(\bar{X}_t)) dt + d\bar{W}_t\end{aligned}$$

for $i = 1, 2, \dots, N$. And $C_{t,i,:} \in \mathbb{R}^n$ denotes the i -th row of the matrix C_t and $D_i(X_t)$ is an $n \times n$ matrix whose elements are all zeros except for the i -th row, which has the vector value $\sigma(w_t X_t + b_t)$.

In the multi-dimensional case, the performance can no longer be measured by the loss function $J(\theta)$, as the numerical computation of the loss J is hard to compute. This difficulty arises during the numerical integration of $\int_0^T (X_t^\theta)^\top Q X_t^\theta + (f_\theta(X_t^\theta))^\top R f_\theta(X_t^\theta) dt$ using numerical discretization. Though for a fixed parameter θ , the time-averaged objective function should converge as $t \rightarrow \infty$. The "Riemann upper sum" of $\int_0^T (X_t^\theta)^\top Q X_t^\theta + (f_\theta(X_t^\theta))^\top R f_\theta(X_t^\theta) dt$ is close to the true value of the integral only when the time step dt of numerical computation is very small. But a

tiny dt will cause the computation really expensive and we have no knowledge about how small the dt should be.

Hence we use the method of error metrics. The error metrics are defined as following:

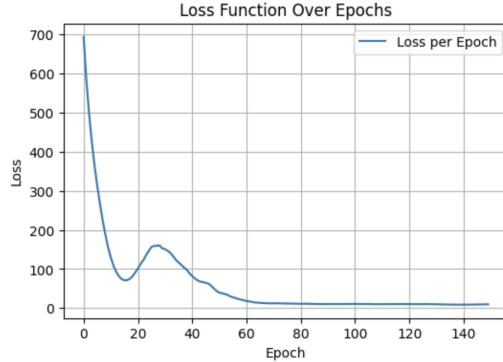
$$\text{Ave Error} = \frac{\sum_{i=1}^n \|f_{\theta_t}(X^i) - \theta^* X^i\|}{\sum_{i=1}^n \|\theta^* X^i\|} \quad (3.17)$$

$$\text{Max Error} = \frac{\max_{i \in \{1, 2, \dots, n\}} \|f_{\theta_t}(X^i) - \theta^* X^i\|}{\sum_{i=1}^n \|\theta^* X^i\|} \quad (3.18)$$

$$\text{Cost Error} = \frac{|J(\theta_T) - J(\theta^*)|}{|C^*|} \quad (3.19)$$

where θ^* is the optimal control and θ_t is the trained parameter. $J(\theta_T)$ and $J(\theta^*)$ denote the objective function $J(\theta)$ in 3.16 with the parameters θ_T and θ^* , respectively.

The numerical results for training the neural network SDE control with the online optimization algorithm are presented in Figure 3.20 and Figure 3.21. In Figure 3.20 we use the loss as the standard while in 3.21 we used error matrices. Overall, the trained neural network control performs effectively, even in high-dimensional settings.



```
Final J: tensor(9.7759, device='cuda:0', grad_fn=<DivBackward0>)
```

Figure 3.20: NN control Multi-D n=5

The code is provided in Appendix A.2

Chapter 4

Futher Research

4.1 Implementation Method Innovation

4.1.1 Methodology Given In The Original Paper

The old way in the original paper is kind of complicated and not well described and heavily depends on the case itself. Hence it is important to give a general law which is applicable in the general case.

4.1.2 An Easier Way

The new way of implementing the algorithm is mainly about using the broadcasting in PyTorch and torch.sum as tools to replace torch.matmul or torch.einsum (einsum notation). For example the update equation in multi-dimensional LQR case:

$$\begin{aligned}\nabla_{\theta} \left[(X_{\ell}^{\theta})^{\top} (Q + \theta^{\top} R \theta) X_{\ell}^{\theta} \right] &= \sum_{i,j} \nabla_{\theta} (q_{i,j} + \theta_{:,i}^{\top} R \theta_{:,j}) X_{\ell}^{\theta,i} X_{\ell}^{\theta,j} \\ &\quad + 2 \sum_{i,j} \nabla_{\theta} X_{\ell}^{\theta,i} (q_{i,j} + \theta_{:,i}^{\top} R \theta_{:,j}) X_{\ell}^{\theta,j} \\ &= \sum_{i,j} ((R\theta)_{k,j} 1_{\{\ell=i\}} + (R\theta)_{k,i} 1_{\{\ell=j\}})_{n \times n} X_{\ell}^{\theta,i} X_{\ell}^{\theta,j} \\ &\quad + 2 \sum_{i,j} \tilde{X}_{\ell}^{\theta,i} (q_{i,j} + \theta_{:,i}^{\top} R \theta_{:,j}) X_{\ell}^{\theta,j},\end{aligned}$$

For the first term $\sum_{i,j} ((R\theta)_{k,j} 1_{\{\ell=i\}} + (R\theta)_{k,i} 1_{\{\ell=j\}})_{n \times n} X_{\ell}^{\theta,i} X_{\ell}^{\theta,j}$. First use for-loop build $\sum_{i,j} ((R\theta)_{k,j} 1_{\{\ell=i\}} + (R\theta)_{k,i} 1_{\{\ell=j\}})_{n \times n}$ which is a n by n matrix.

```
D = torch.zeros(n, n, n, n, device=device)
for k in range(dim):
    for l in range(dim):
```

```

D[k, 1, :, k] += Rtheta[:, 1]
D[k, 1, :, 1] += Rtheta[:, k]

```

And notice here $X_t^{\theta,i}$ and $X_t^{\theta,j}$ are n dimensional tensor, i.e. without the batch dimension. Since it is obtained after the torch.mean on the batch dimension. Our thoughts can be conclude as the equation below:

$$((R\theta)_{:,j} \mathbf{1}_{\{i=n\}} + (R\theta)_{:,i} \mathbf{1}_{\{j=n\}}) \rightarrow (n, n) \xrightarrow{\text{.unsqueeze}()} (1, 1, n, n) = M \quad (4.1)$$

$$X_t^{\theta,i} \rightarrow (n) \xrightarrow{\text{.unsqueeze}()} (n, 1, 1, 1) = X_1 \quad (4.2)$$

$$X_t^{\theta,j} \rightarrow (n) \xrightarrow{\text{.unsqueeze}()} (1, n, 1, 1) = X_2 \quad (4.3)$$

where the index of the 4 dimensional tensors is (i, j, k, l) .

Then with the tool broadcast in PyTorch, we get $M * X_1 * X_2$, the final step is to sum with respect to dimension i and j, which is torch.sum(torch.sum(, dim=1), dim=0).

For the second term $2 \sum_{i,j} \tilde{X}_t^{\theta,i} (q_{i,j} + \theta_{:,i}^\top R\theta_{:,j}) X_l^{\theta,j}$. The implement idea is described below:

$$\tilde{X}_t^\theta \rightarrow (n, n, n) \xrightarrow{\text{.unsqueeze}()} (n, 1, n, n) = X_{tilde} \quad (4.4)$$

$$(q_{i,j} + \theta_{:,i}^\top R\theta_{:,j}) \rightarrow (n, n) \xrightarrow{\text{.unsqueeze}()} (n, n, 1, 1) = Middle \quad (4.5)$$

$$X_t^\theta \rightarrow (n) \xrightarrow{\text{.unsqueeze}()} (1, n, 1, 1) = X_3 \quad (4.6)$$

where the index of the 4-d tensor is (i, j, k, l) .

Then with the tool broadcast in PyTorch, we get $X_{tilde} * Middle * X_3$, the final step is to sum with respect to dimension i and j, which is torch.sum(torch.sum(, dim=1), dim=0). Mind the 2 in the front of the 2nd term.

The above discussions shows the main idea of this new implementation method. With the idea in mind, we can deal with more complicated case. Another example is also in the update equation in multi-dimensional LQR case:

$$d\tilde{X}_t^{\theta,i} = \left(-\tilde{X}_t^{\theta,i} + \sum_{j=1}^n \theta_{i,j} \tilde{X}_t^{\theta,j} + D_i(X_t^\theta) \right) dt \quad (4.7)$$

where $D_i(X_t^\theta)$ is an $n \times n$ matrix whose elements are all zeros except for the i -th row, which has values X_t^θ .

To deal with $\sum_{j=1}^n \theta_{i,j} \tilde{X}_t^{\theta,j}$ same thing is "repeated" below:

$$\tilde{X}_t^\theta \rightarrow (N, n, n, n) \xrightarrow{\text{unsqueeze}} (N, 1, n, n, n) \quad (4.8)$$

$$\theta \rightarrow (n, n) \xrightarrow{\text{unsqueeze}} (1, n, n, 1, 1) \quad (4.9)$$

where the 5-d tensor is indexed by $(batch, i, j, a, b)$.

Then we sum with respect to dimension j , and get a 4-d tensor $(batch, i, a, b)$.

To avoid using for-loops, the first question here is how to build D as a hyper-cube shown below. Our method is:

1. Our idea is to first produce an eye(n) matrix in dimension (i, b) .
2. Then unsqueeze it to get $(1, n, 1, n)$ with index $(batch, i, a, b)$.
3. broadcast the above 4-d tensor with unsqueezed X_t^θ , which has 4-dimensions $(N, 1, n, 1)$ with index $(batch, i, a, b)$.

This provide us with the hyper-cube D.

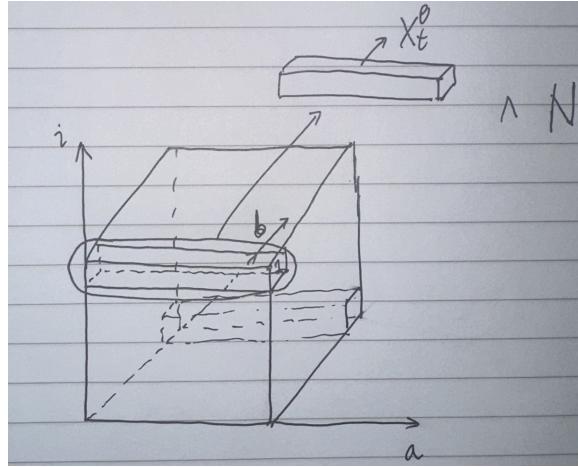


Figure 4.1: Hyper Cube

The final step of the implementation is to sum the $-X_t^\theta$, $\sum_{j=1}^n \theta_{i,j} \tilde{X}_t^{\theta,j}$ and D .

4.1.3 Advantages

Using the new implementation idea, we can deal with multiplications between hyper-matrix as well as multiplications between tensors with different dimensions. The main advantages are as follows. The new method is universal in the sense that it applies no matter what dimensions the tensors are. And the user does not need

to remember the names and signatures of all the different functions in PyTorch for calculating dot products, outer products, transposes and matrix-vector or matrix-matrix multiplications. And the method is so natural that once you understand the equations you can implement those equations using the given idea.

4.2 Higher Dimensions

4.2.1 LQR

Since the higher the dimension is the better GPU with higher RAM to train the model. Even H100 (which is the best) only got 80GB RAM, hence the only way is to lower the batch size, which will make the output noisy. This is a trade off. After taking those factors into account, we try the case which dimension equals to $dim = 50$.

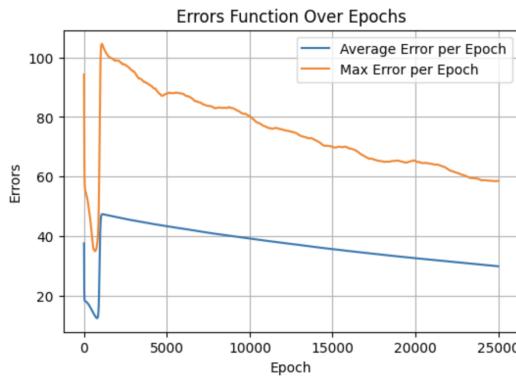


Figure 4.2: LQR n=50

Illustrating the changes in average and maximum errors across 25,000 epochs during a 4-hour runtime, the plot indicates that while the model has achieved a considerable reduction in average error, resulting in overall improved performance, the presence of slow convergence suggests a larger batch size and a longer run time will be beneficial. It is worth mentioning there is a trade-off between run time and computational cost involved. With sufficient computation resources such as GPU RAM and runtime, the algorithm can perform well.

4.2.2 Neural Network Control

After taking factors (RAM, batch size, dimensions) into account, we try the case which dimension equals to $dim = 30$. The plot indicates that while the online SGD

algorithm effectively reduces the loss initially, it faces challenges in achieving stable convergence, highlighting areas for potential hyper parameter tuning and optimization strategy adjustments.

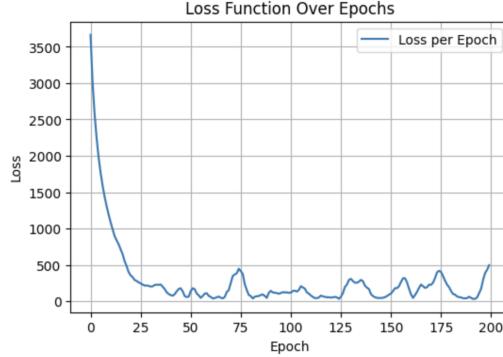


Figure 4.3: NN n=30

Note that while using the loss as the metric to measure performance, we need large RAM size to perform high dimensional Neural Network examples.

4.2.3 Conclusions

Though in very high dimension cases the algorithm requires a lot of computational resources, it does provide a feasible way to solve high dimensional control problems, and the computational costs are acceptable.

4.3 Correlated OU Processes

For the parameters $\theta = (\mu, \sigma)$ with $\mu \in \mathbb{R}^m$, $\sigma \in \mathbb{R}^{m \times m}$, let the m -dimensional process X_t^θ satisfy

$$dX_t^\theta = (\mu - X_t^\theta) dt + \sigma dW_t \quad (4.10)$$

where $W_t \in \mathbb{R}^m$.

In the section 3.1.4 Multi-Dimensional Ornstein-Uhlenbeck Process, we pick σ to be an identity matrix or a matrix similar to an identity matrix. Now we try to add some more general correlation into the case using the following method.

```
# Create an identity matrix
identity_matrix = torch.eye(m, m, device='cuda')
# Add small random perturbations
```

```

perturbations = torch.randn(m, m, device='cuda') * 0.01
# Sigma
sigma = identity_matrix + perturbations

```

now the algorithm is the same as in section 3.1.4:

$$\begin{aligned}
d\mu_t &= -4\alpha_t \left(|\bar{X}_t|^2 - 2m \right) \left(\sum_{k=1}^m X_t^k \tilde{X}_t^{\mu,k} \right) dt \\
d\sigma_t &= -4\alpha_t \left(|\bar{X}_t|^2 - 2m \right) \left(\sum_{k=1}^m X_t^k \tilde{X}_t^{\sigma,k} \right) dt \\
dX_t &= (\mu_t - X_t) dt + \sigma_t dW_t \\
d\bar{X}_t &= (\mu_t - \bar{X}_t) dt + \sigma_t d\bar{W}_t \\
d\tilde{X}_t^\mu &= (I_m - \tilde{X}_t^\mu) dt \\
d\tilde{X}_t^{\sigma,i} &= -\tilde{X}_t^{\sigma,i} dt + D_i(dW_t), \quad i \in \{1, \dots, m\}
\end{aligned}$$

where I_m is the $m \times m$ identity matrix and where $D_i(dW_t)$ is a $m \times m$ matrix with all elements equal to 0 except i -th column being dW_t .

For dimensions $m = 10$. In Figures 4.4, we observe though the algorithm converge the performance is not so good. But it is worth mentioning that for some other σ matrix, as in Figure 4.5, the loss converges to a smaller value. Hence we reckon that for correlated cases, the correlation can influence the performance.

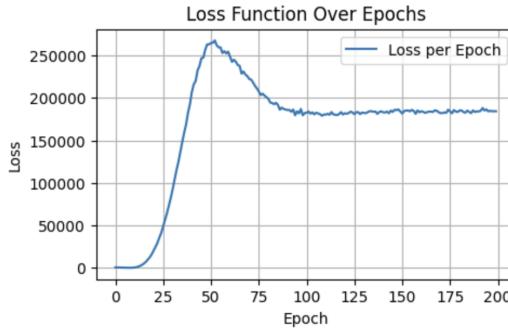


Figure 4.4: Correlated dim=10 (1)

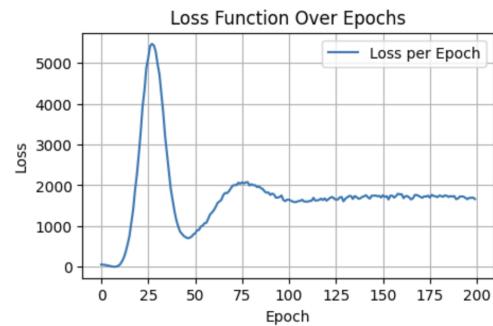


Figure 4.5: Correlated dim=10 (2)

Chapter 5

Conclusions

5.1 Summary of Main Findings

This dissertation presents a comprehensive study on the online optimization of neural network controls for stochastic differential equations. Through the re-implementation of existing methods and exploration of new dimensions and correlations, the research validates the proposed algorithm's capability to handle high-dimensional stochastic optimal control problems. The continuous-time stochastic gradient descent method proved effective in various scenarios, including linear and nonlinear processes, path-dependent SDEs, and multi-dimensional cases. Despite certain theoretical limitations such as dependency on hyper parameters and noise in stochastic estimates, the algorithm shows promise for practical applications. The findings underscore the importance of careful parameter tuning and highlight the potential for significant advancements in optimizing neural network controls for high-dimensional SDEs.

5.2 Future Research

Exploring correlative cases and non-linear processes with the given algorithm opens new opportunities. Correlative cases involve interdependent stochastic processes, introducing complexities that broaden the algorithm's applicability. Investigating how the algorithm handles correlations between state variables and control processes is crucial for real-world implementation.

Non-linear processes pose challenges due to their dynamic state evolution, requiring thorough examination of the algorithm's stability, convergence, and computational feasibility. Extending the algorithm to non-linear processes would showcase its robustness and versatility for various applications, including financial markets and engineering systems.

Additionally, a rigorous mathematical proof of the algorithm's effectiveness in general cases would provide a solid theoretical foundation. Proving stability, convergence, and establishing error bounds under general conditions ensures that empirical performance is supported by robust principles.

In summary, extending the algorithm to handle correlative cases and non-linear processes is a promising research direction. Demonstrating its capabilities in these contexts and providing mathematical proof would enhance its theoretical foundation and expand its practical applicability.

Appendix A

Key Code

We listed pieces of code below, those are the code for the complicated cases in the dissertation.

A.1 One-dimensional Neural Network Control

```
import torch
import matplotlib.pyplot as plt
from tqdm import tqdm

# Check if CUDA is available and set the device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
# device = 'cpu'
# Hyperparameters
N = 1000 # Batch size
I = 100 # Total number of iterations
dt = 0.05 # Time increment
m = 10

# Initialization
# Example vectors for c, w, b
c = torch.full(size=(m,), fill_value=1, device=device).float() # vector c
w = torch.full(size=(m,), fill_value=4, device=device).float() # vector w
b = torch.full(size=(m,), fill_value=6, device=device).float() # vector b
# Stack the vectors into a matrix where each vector is a row
theta = torch.stack((w, b, c))
```

```

# Initialization for X and its tilde variables
X = torch.ones(N, requires_grad=True, device=device)
X_bar = torch.ones(N, requires_grad=True, device=device)
X_tilde_w = torch.zeros(N, m, device=device)
X_tilde_b = torch.zeros(N, m, device=device)
X_tilde_c = torch.zeros(N, m, device=device)

print("Finished with initialization.")

def sigmoid_derivative(x):
    return torch.sigmoid(x) * (1 - torch.sigmoid(x))

# Create class for convience
class SingleLayerNNManual:
    def __init__(self, theta):
        super(SingleLayerNNManual, self).__init__()
        # Manually initialize weights and biases
        self.weights = theta[0, :].float().clone().requires_grad_(True)
        self.biases = theta[1, :].float().clone().requires_grad_(True)
        self.output_weights = theta[2, :].float().clone().requires_grad_(True)

    def forward(self, x):
        # Define the function
        f = torch.sum(self.output_weights * torch.sigmoid(self.weights *
            x.unsqueeze(1) + self.biases), dim = 1)
        # Return
        return f
    ,

    def zero_grad(self):
        # Manually zero the gradients
        if self.weights.grad is not None:
            self.weights.grad.zero_()
        if self.biases.grad is not None:

```

```

        self.biases.grad.zero_()
    if self.output_weights.grad is not None:
        self.output_weights.grad.zero_()
    ...

# Create class
class OUSimulator:
    def __init__(self, model, device, batch_size=1000, dt=0.1, num_steps=1000):
        self.model = model
        self.device = device
        self.batch_size = batch_size
        self.dt = dt
        self.num_steps = num_steps
        self.X = torch.ones(self.batch_size, device=self.device)
        self.integral_sum = 0.0

    def simulate(self):
        """Run the Ornstein-Uhlenbeck process simulation."""
        for i in range(self.num_steps):
            dW = torch.randn(self.batch_size, device=self.device) *
                torch.sqrt(torch.tensor(self.dt, device=self.device))
            self.X += (-self.X + self.model.forward(self.X)) * self.dt + dW
            self.integral_sum += torch.mean(self.X**2) *
                torch.mean(self.model.forward(self.X)**2) * self.dt

    def compute_loss(self):
        """Calculate the loss function J based on the simulation."""
        loss = self.integral_sum / (self.dt * self.num_steps)
        return loss.cpu().item()

# The containers
loss_history = []
theta_history = []
gradient_difference = []

```

```

# Loop with defined range
for t in tqdm(range(I), "Loop within defined range..."):
    # Learning rate adjustment
    if t < 50:
        alpha = 0.07
    elif t < 75:
        alpha = 0.03
    else:
        alpha = 0.01

    # Create an instance of the network
    model = SingleLayerNNManual(theta)
    output = model.forward(X)

    gradient = torch.sum(c * sigmoid_derivative(w * torch.mean(X) + b) * w)

    # output.sum().backward()
    # gradient = X.grad
    # print(gradient)

    # Update the changes
    theta[0, :] = theta[0, :] + (-alpha) * (2*torch.mean(X)*torch.mean(X_tilde_w,
dim=0) + 2 * torch.mean(output) *
(c*sigmoid_derivative(w*torch.mean(X)+b)*torch.mean(X) + torch.mean(gradient) *
torch.mean(X_tilde_w, dim=0))) * dt
    theta[1, :] = theta[1, :] + (-alpha) * (2*torch.mean(X)*torch.mean(X_tilde_b,
dim=0) + 2 * torch.mean(output) * (c*sigmoid_derivative(w*torch.mean(X)+b) +
torch.mean(gradient) * torch.mean(X_tilde_b, dim=0))) * dt
    theta[2, :] = theta[2, :] + (-alpha) * (2*torch.mean(X)*torch.mean(X_tilde_c,
dim=0) + 2 * torch.mean(output) * (torch.sigmoid(w*torch.mean(X)+b) +
torch.mean(gradient) * torch.mean(X_tilde_c, dim=0))) * dt

    # BM increment
    dW = torch.randn(N, device=device) * torch.sqrt(torch.tensor(dt,
device=device))

```

```

dW_bar = torch.randn(N, device=device) * torch.sqrt(torch.tensor(dt,
device=device))

# The changes
X = X + (-X + output) * dt + dW
X_tilde_w = X_tilde_w + (-X_tilde_w + c.repeat(N, 1) *
sigmoid_derivative(torch.matmul(X.unsqueeze(1), w.unsqueeze(0)) + b.repeat(N,
1)) * torch.mean(X) + torch.mean(gradient) * X_tilde_w) * dt
X_tilde_b = X_tilde_b + (-X_tilde_b + c.repeat(N, 1) *
sigmoid_derivative(torch.matmul(X.unsqueeze(1), w.unsqueeze(0)) + b.repeat(N,
1)) + torch.mean(gradient) * X_tilde_b) * dt
X_tilde_c = X_tilde_c + (-X_tilde_c +
torch.sigmoid(torch.matmul(X.unsqueeze(1), w.unsqueeze(0)) + b.repeat(N, 1)) +
torch.mean(gradient) * X_tilde_c) * dt
X_bar = X_bar + (-X_bar + model.forward(X_bar)) * dt + dW_bar

# Store theta
theta_history.append(theta.cpu().detach().numpy())

# Simulate the OU
model_simu = SingleLayerNNManual(theta) # Assuming SingleLayerNNManual is
defined elsewhere

simulator = OUSimulator(model=model_simu, device=device)
simulator.simulate()
loss = simulator.compute_loss()
loss_history.append(loss)

# Check the gradient
eps = 1e-5

w = theta[0, :]
w0 = w[0]
autograd_gradient_w = (2*torch.mean(X)*torch.mean(X_tilde_w, dim=0) + 2 *
torch.mean(output) * (c*sigmoid_derivative(w*torch.mean(X)+b)*torch.mean(X) +
torch.mean(gradient) * torch.mean(X_tilde_w, dim=0))) * dt

```

```

autograd_gradient_w0 = autograd_gradient_w[0]

x_perturb_plus = w0.clone()
x_perturb_minus = w0.clone()
# plus
x_perturb_plus += eps
theta[0, 0] = x_perturb_plus
model_check = SingleLayerNNManual(theta) # Assuming SingleLayerNNManual is
defined elsewhere
simulator = OUSimulator(model=model_check, device=device)
simulator.simulate()
loss_plus = simulator.compute_loss()

# minus
x_perturb_minus -= eps
theta[0, 0] = x_perturb_minus
model_check = SingleLayerNNManual(theta) # Assuming SingleLayerNNManual is
defined elsewhere
simulator = OUSimulator(model=model_check, device=device)
simulator.simulate()
loss_minus = simulator.compute_loss()

grad_by_fd = (loss_plus - loss_minus) / (2 * eps)

del x_perturb_minus
del x_perturb_plus

# Compare gradients
gradient_difference.append(torch.norm(autograd_gradient_w0 -
grad_by_fd).cpu().item())
print(f"Gradient difference at iteration {t}:
{torch.norm(autograd_gradient_w0 - grad_by_fd).item()}")


# Zero grad
# model.zero_grad()

```

```

del X, X_bar, X_tilde_w, X_tilde_b, X_tilde_c # Delete tensors that are no
longer needed
torch.cuda.empty_cache() # Clear cache

print("Start visualization.")
# Plotting the check results over epochs
plt.figure(figsize=(6, 4))
plt.plot(gradient_difference, label='Gd per Epoch')
plt.title('Gd Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Gd')
plt.legend()
plt.grid(True)
plt.show()

# Plotting the loss over epochs
plt.figure(figsize=(6, 4))
plt.plot(loss_history, label='Loss per Epoch')
plt.title('Loss Function Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()

# Output the final values of parameters
print("Final theta:", theta)
print("Final J:", loss)

```

A.2 Multi-dimensional Neural Network Control

```

import torch
import matplotlib.pyplot as plt

```

```

from tqdm import tqdm
import torch.nn as nn
import torch.optim as optim

# Check if CUDA is available and set the device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Hyperparameters
N = 100 # Batch size
I = 200 # Total number of iterations
dt = 0.05 # Time increment
n = 30 # Dimension
m = 10

# Example Q and R matrices
Q = torch.eye(n, device=device).double()
R = torch.eye(n, device=device).double()

# Example vectors for c, w, b
c = torch.randn(n, m, device=device).double()
w = torch.randn(m, n, device=device).double()
b = torch.randn(m, device=device).double()

# Initialization for X and its tilde variables
X = torch.ones(N, n, requires_grad=False, device=device).double()
X_0 = torch.ones(N, n, requires_grad=False, device=device).double()
X_0_J = torch.ones(5000, n, requires_grad=False, device=device).double() #
Reduced batch size for X_0_J
X_bar = torch.ones(N, n, requires_grad=False, device=device).double()
X_tilde_w = torch.zeros(N, n, m, n, device=device).double()
X_tilde_b = torch.zeros(N, n, m, device=device).double()
X_tilde_c = torch.zeros(N, n, n, m, device=device).double()

# Containers
loss_history = []

```

```

gradient_difference = []

print("Finished with initialization.")

# Functions and classes
def sigmoid_derivative(x):
    return torch.sigmoid(x) * (1 - torch.sigmoid(x))

class SingleLayerNN(nn.Module):
    def __init__(self, w, b, c):
        super(SingleLayerNN, self).__init__()
        self.w = nn.Parameter(w)
        self.b = nn.Parameter(b)
        self.c = nn.Parameter(c)
        self.activation = nn.Sigmoid()

    def forward(self, x):
        hidden = self.activation(torch.matmul(self.w, x) + self.b)
        output = torch.matmul(self.c, hidden)
        return output

class SingleLayerNNBatch(nn.Module):
    def __init__(self, w, b, c, N = 100, m = 10):
        super(SingleLayerNNBatch, self).__init__()
        self.w = nn.Parameter(w)
        self.b = nn.Parameter(b)
        self.c = nn.Parameter(c)
        self.N = N
        self.m = m
        self.activation = nn.Sigmoid()

    def forward(self, x):
        # x is of shape (batch_size, input_dim)
        hidden = self.activation(torch.einsum('ij, bj -> bi', self.w, x) +
                               self.b.unsqueeze(0).expand(self.N, self.m))
        output = torch.einsum('ij, bj -> bi', self.c, hidden)

```

```

        return output

class Simulator:
    def __init__(self, model_1, model_Batch, n, m, Q, R, dt=0.01, T=100.0,
                 timesteps=10000.0, batch_size=5000): # Reduced batch size
        self.batch_size = batch_size
        self.model_1 = model_1
        self.model_Batch = model_Batch
        self.input_dim = n
        self.hidden_dim = m
        self.output_dim = n
        self.Q = Q
        self.R = R
        self.dt = dt
        self.T = T
        self.timesteps = 10000

    def simulate_X_t(self, X_0):
        X_0_mean = torch.mean(X_0, dim=0)
        X_t = X_0
        X_t_trajectory = [X_0_mean]

        for _ in range(self.timesteps):
            dW_t = torch.randn(self.batch_size, self.input_dim,
                               device=device).double() * torch.sqrt(torch.tensor(self.dt).double())
            f_theta_X_t = self.model_Batch.forward(X_t)
            X_t = X_t + (-X_t + f_theta_X_t) * self.dt + dW_t
            X_t_mean = torch.mean(X_t, dim=0)
            X_t_trajectory.append(X_t_mean)

        return torch.stack(X_t_trajectory)

    def compute_loss(self, X_trajectory):
        integral = 0

        for X_t_mean in X_trajectory:

```

```

        f_theta_X_t_mean = self.model_1.forward(X_t_mean)
        term1 = torch.matmul(X_t_mean.t(), torch.matmul(self.Q, X_t_mean))
        term2 = torch.matmul(f_theta_X_t_mean.t(), torch.matmul(self.R,
        f_theta_X_t_mean))
        integral += term1 + term2

    J_theta = integral / self.T
    return J_theta

def forward(self, X_0):
    X_trajectory = self.simulate_X_t(X_0)
    loss = self.compute_loss(X_trajectory)
    return loss

# tqdm for loop
for t in tqdm(range(I), "Loop within defined range..."):
    # Learning rate adjustment
    if t < 50:
        alpha = 0.2
    elif t < 90:
        alpha = 0.2
    else:
        alpha = 0.1

    # Simulate the loss J(theta)
    model_simu_1 = SingleLayerNN(w, b, c).to(device)
    model_simu_batch = SingleLayerNNBatch(w, b, c, 5000).to(device) # Reduced
    batch size
    simulator = Simulator(model_simu_1, model_simu_batch, n, m, Q, R)
    loss = simulator.forward(X_0_J)
    print('loss', loss)
    loss_history.append(loss.cpu().item())

    # Compute the mean
    X_mean_o = torch.mean(X, dim=0)
    X_mean = X_mean_o.clone().detach().requires_grad_(True)

```

```

X_tilde_w_mean = torch.mean(X_tilde_w, dim=0)
X_tilde_b_mean = torch.mean(X_tilde_b, dim=0)
X_tilde_c_mean = torch.mean(X_tilde_c, dim=0)

# Create an instance of the network
model = SingleLayerNN(w, b, c).to(device)
f_theta = model(X_mean)
model_batch = SingleLayerNNBatch(w, b, c).to(device)
f_theta_batch = model_batch(X)

# Compute Gradient w.r.t w&b&c
first = torch.matmul(f_theta.t(), torch.matmul(R, f_theta))
term1 = torch.matmul(X_mean.t(), torch.matmul(Q, X_mean))
term2 = torch.matmul(f_theta.t(), torch.matmul(R, f_theta))
second = term1 + term2

# Compute gradients with respect to X_t
second.backward(retain_graph=True)

# Compute the additional term for parameter updates
grad_X_t = X_mean.grad
sum_term_w = torch.sum(grad_X_t.unsqueeze(1).unsqueeze(2) * X_tilde_w_mean,
dim=0)
sum_term_b = torch.sum(grad_X_t.unsqueeze(1) * X_tilde_b_mean, dim=0)
sum_term_c = torch.sum(grad_X_t.unsqueeze(1).unsqueeze(2) * X_tilde_c_mean,
dim=0)

# Compute the gradient
first.backward()

# Print gradients
first_w = model.w.grad
first_b = model.b.grad
first_c = model.c.grad

# Update the changes

```

```

dw = (-alpha) * (first_w + sum_term_w) * dt
db = (-alpha) * (first_b + sum_term_b) * dt
dc = (-alpha) * (first_c + sum_term_c) * dt

# BM increment
dW = torch.randn_like(X) * torch.sqrt(torch.tensor(dt).double()).to(device)

# The changes
dX = (-X + f_theta_batch) * dt + dW
wXb = torch.einsum('ij, bj -> bi', w, X) + b.unsqueeze(0).expand(N, m)
sigmoid_dash = sigmoid_derivative(wXb)
sigmoidwXb = torch.sigmoid(wXb)
dX_tilde_w = torch.zeros(N, n, m, n, device=device).double()
dX_tilde_b = torch.zeros(N, n, m, device=device).double()
dX_tilde_c = torch.zeros(N, n, n, m, device=device).double()
for i in range(n):
    Di = torch.zeros(N, n, m, device=device).double()
    Di[:, i, :] = sigmoidwXb
    dX_tilde_w[:, i, :, :] = (-X_tilde_w[:, i, :, :] + torch.sum(c[i, :].unsqueeze(0).unsqueeze(2).unsqueeze(3) *
        sigmoid_dash.unsqueeze(2).unsqueeze(3) *
        torch.sum(w.unsqueeze(0).unsqueeze(3).unsqueeze(4) *
        X_tilde_w.unsqueeze(1), dim=2), dim=1) + torch.bmm((c[i, :] * sigmoid_dash).unsqueeze(2), X.unsqueeze(1))) * dt
    dX_tilde_b[:, i, :] = (-X_tilde_b[:, i, :] + torch.sum(c[i, :].unsqueeze(0).unsqueeze(2) * sigmoid_dash.unsqueeze(2) *
        torch.sum(w.unsqueeze(0).unsqueeze(3) * X_tilde_b.unsqueeze(1), dim=2),
        dim=1) + c[i, :] * sigmoid_dash) * dt
    dX_tilde_c[:, i, :, :] = (-X_tilde_c[:, i, :, :] + torch.sum(c[i, :].unsqueeze(0).unsqueeze(2).unsqueeze(3) *
        sigmoid_dash.unsqueeze(2).unsqueeze(3) *
        torch.sum(w.unsqueeze(0).unsqueeze(3).unsqueeze(4) *
        X_tilde_c.unsqueeze(1), dim=2), dim=1) + Di) * dt

# Update
w = w + dw

```

```

    b = b + db
    c = c + dc
    X = X + dX
    X_tilde_w = X_tilde_w + dX_tilde_w
    X_tilde_b = X_tilde_b + dX_tilde_b
    X_tilde_c = X_tilde_c + dX_tilde_c

print("Start visualization.")

# Plotting the loss over epochs
plt.figure(figsize=(6, 4))
plt.plot(loss_history, label='Loss per Epoch')
plt.title('Loss Function Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()

# Output the final loss
print("Final J:", loss)

```

A.3 LQR

```

import torch
import matplotlib.pyplot as plt
from tqdm import tqdm

# Check if CUDA is available and set the device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
# device = 'cpu'

# f(Y) = Y^2

# Hyperparameters
N = 10000 # Batch size

```

```

I = 100 # Total number of iterations
dt = 0.05 # Time increment

# Initialization
theta = torch.tensor(1, device=device, dtype=torch.float32)

# Initialization for X and its tilde variables
X = torch.ones(N, device=device)
X_tilde = torch.zeros(N, device=device)

print("Finished with initialization.")

# A container for the loss J
loss_history = []
theta_history = []

# Generate the Brownian increament
# increments = torch.randn(size=(2*N, I), device=device) *
torch.sqrt(torch.tensor(dt))

# Loop with defined range
for t in tqdm(range(I), "Loop with defined range..."):
    # Learning rate adjustment
    if t < 50:
        alpha = 2
    elif t < 75:
        alpha = 1
    else:
        alpha = 0.5

    # Calculate the changes
    dtheta = (-alpha) * torch.mean(2 * theta * X**2 + 2 * (1 + theta**2) * X *
X_tilde) * dt

    # BM increment
    dW = torch.randn(N, device=device) * torch.sqrt(torch.tensor(dt,

```

```

device=device))

dW_bar = torch.randn(N, device=device) * torch.sqrt(torch.tensor(dt,
device=device))

# Make sure that t is zero-based (i.e., first column is i=0)
#dW = increaments[:N, t] # t_th_column
#dW_bar = increaments[N:, t]

# The changes
dX = (theta - 1) * X * dt + dW
dX_tilde = (X + (theta - 1) * X_tilde) * dt

# Update parameters
theta += dtheta
X += dX
X_tilde += dX_tilde

# Store theta
theta_history.append(theta.cpu().item())

# Simulate the OU process
batch_size_simu = 10000
dt_simu = 0.1
number_of_step_simu = 10000
X_simu = torch.ones(batch_size_simu, device=device)
integral_sum = 0
# increaments_simu = torch.randn(size=(batch_size_simu, number_of_step_simu),
device=device) * torch.sqrt(torch.tensor(dt_simu, device=device))
for i in range(0, number_of_step_simu):
    dW = torch.randn(batch_size_simu, device=device) *
    torch.sqrt(torch.tensor(dt_simu, device=device))
    #dw = increaments_simu[:,i]
    X_simu = X_simu + (-X_simu + theta * X_simu) * dt_simu + dW
    integral_sum += (1 + theta**2) * torch.mean(X_simu**2) * dt

# Compute the loss function J
loss = integral_sum / (dt_simu * number_of_step_simu)

```

```

loss_history.append(loss.cpu().item())

print("Start visualization.")

# Plotting the loss over epochs
plt.figure(figsize=(10, 6))
plt.plot(loss_history, label='Loss per Epoch')
plt.title('Loss Function Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()

# Plotting the theta over epochs
plt.figure(figsize=(10, 6))
plt.plot(theta_history, label='Theta per Epoch')
plt.title('Theta Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Theta')
plt.legend()
plt.grid(True)
plt.show()

# Output the final values of parameters
print("Final theta1:", theta.item())
print("Final J:", loss)

```

Bibliography

- [1] Brian DO Anderson and John B Moore. *Optimal control: linear quadratic methods*. Courier Corporation, 2007.
- [2] Michael Röckner, Xiaobin Sun, and Yingchao Xie. Strong convergence order for slow-fast mckean-vlasov stochastic differential equations, 2019.
- [3] Ziheng Wang and Justin Sirignano. A forward propagation algorithm for online optimization of nonlinear stochastic differential equations, 2022.
- [4] Ziheng Wang and Justin Sirignano. Continuous-time stochastic gradient descent for optimizing over the stationary distribution of stochastic differential equations, 2023.